


Hoofdstuk 14

Externe Opslag: Basisconcepten

Wat is data? Bits : 2 \neq voltages



(a) Logic diagram

S	R	Q	Q'
1	0	1	0
0	0	1	0
0	1	0	1
0	0	0	1
1	1	0	0

(after S=1, R=0)


(after S=0, R=1)

(b) Truth table

Basic flip-flop circuit with NOR gates

“binary digit”

Wat is data? Bytes



Grootste byte: 255

256 ≠ waarden

Natuurlijke getallen als byterijen

```
> (define number 1234567)
> (modulo number 256)
135
> (set! number (quotient number 256))
> (modulo number 256)
214
> (set! number (quotient number 256))
> number
18
```

$$1234567 = 18 \cdot 256^2 + 214 \cdot 256^1 + 135 \cdot 256^0$$

18, 214, 135

big endian


meest significante byte op
eerste (laagste)
geheugenadres

135, 214, 18

little endian

minst significante byte op
eerste (laagste)
geheugenadres

Grafische data als byterijen



60, 66, 2, 30, 98, 66, 70, 59

Terminologie

1024 bytes = 1 kilobyte $1024 = 2^{10}$
1024 kilobyte = 1 megabyte (1M)
1024 megabyte = 1 gigabyte (1G)
1024 gigabyte = 1 terabyte (1T)
1024 terabyte = 1 petabyte (1P)
1024 petabyte = 1 exabyte (1E)
1024 exabyte = 1 zettabyte (1Z)
1024 zettabyte = 1 yottabyte (1Y)

Computers
 \approx
gigabytes


Harddisks
 \approx
terabytes

$2^0 = 1$
 $2^1 = 2$
 $2^2 = 4$
 $2^3 = 8$
 $2^4 = 16$
 $2^5 = 32$
 $2^6 = 64$
 $2^7 = 128$
 $2^8 = 256$
 $2^{10} = 1024$
 $2^{16} = 65536$

Scheme's Bytevectoren

Nuttig om "rauwe"
data voor te stellen

a



00111100
01000010
00000010
00011110
01100010
01000010
01000110
00111011

```
(define letter-a  
  (let ((v (make-bytevector 8)))  
    (bytevector-u8-set! v 0 60)  
    (bytevector-u8-set! v 1 66)  
    (bytevector-u8-set! v 2 2)  
    (bytevector-u8-set! v 3 30)  
    (bytevector-u8-set! v 4 98)  
    (bytevector-u8-set! v 5 66)  
    (bytevector-u8-set! v 6 70)  
    (bytevector-u8-set! v 7 59)  
    v))
```

Veel efficiënter dan
gewone vectoren

values + type tag

```
> (bytevector-u8-ref letter-a 6)  
70
```

make-bytevector, bytevector-u8-ref,
bytevector-u8-set!

∈ R7RS

Getallen in Bytevectors: plaats

```
(define (natural-bytes nmr)  
  (exact (ceiling (log (max (+ nmr 1) 2) 256))))
```

```
(define (integer-bytes nmr)  
  (exact (ceiling (log (max (abs (+ (* 2 nmr) 1)) 2) 256))))
```

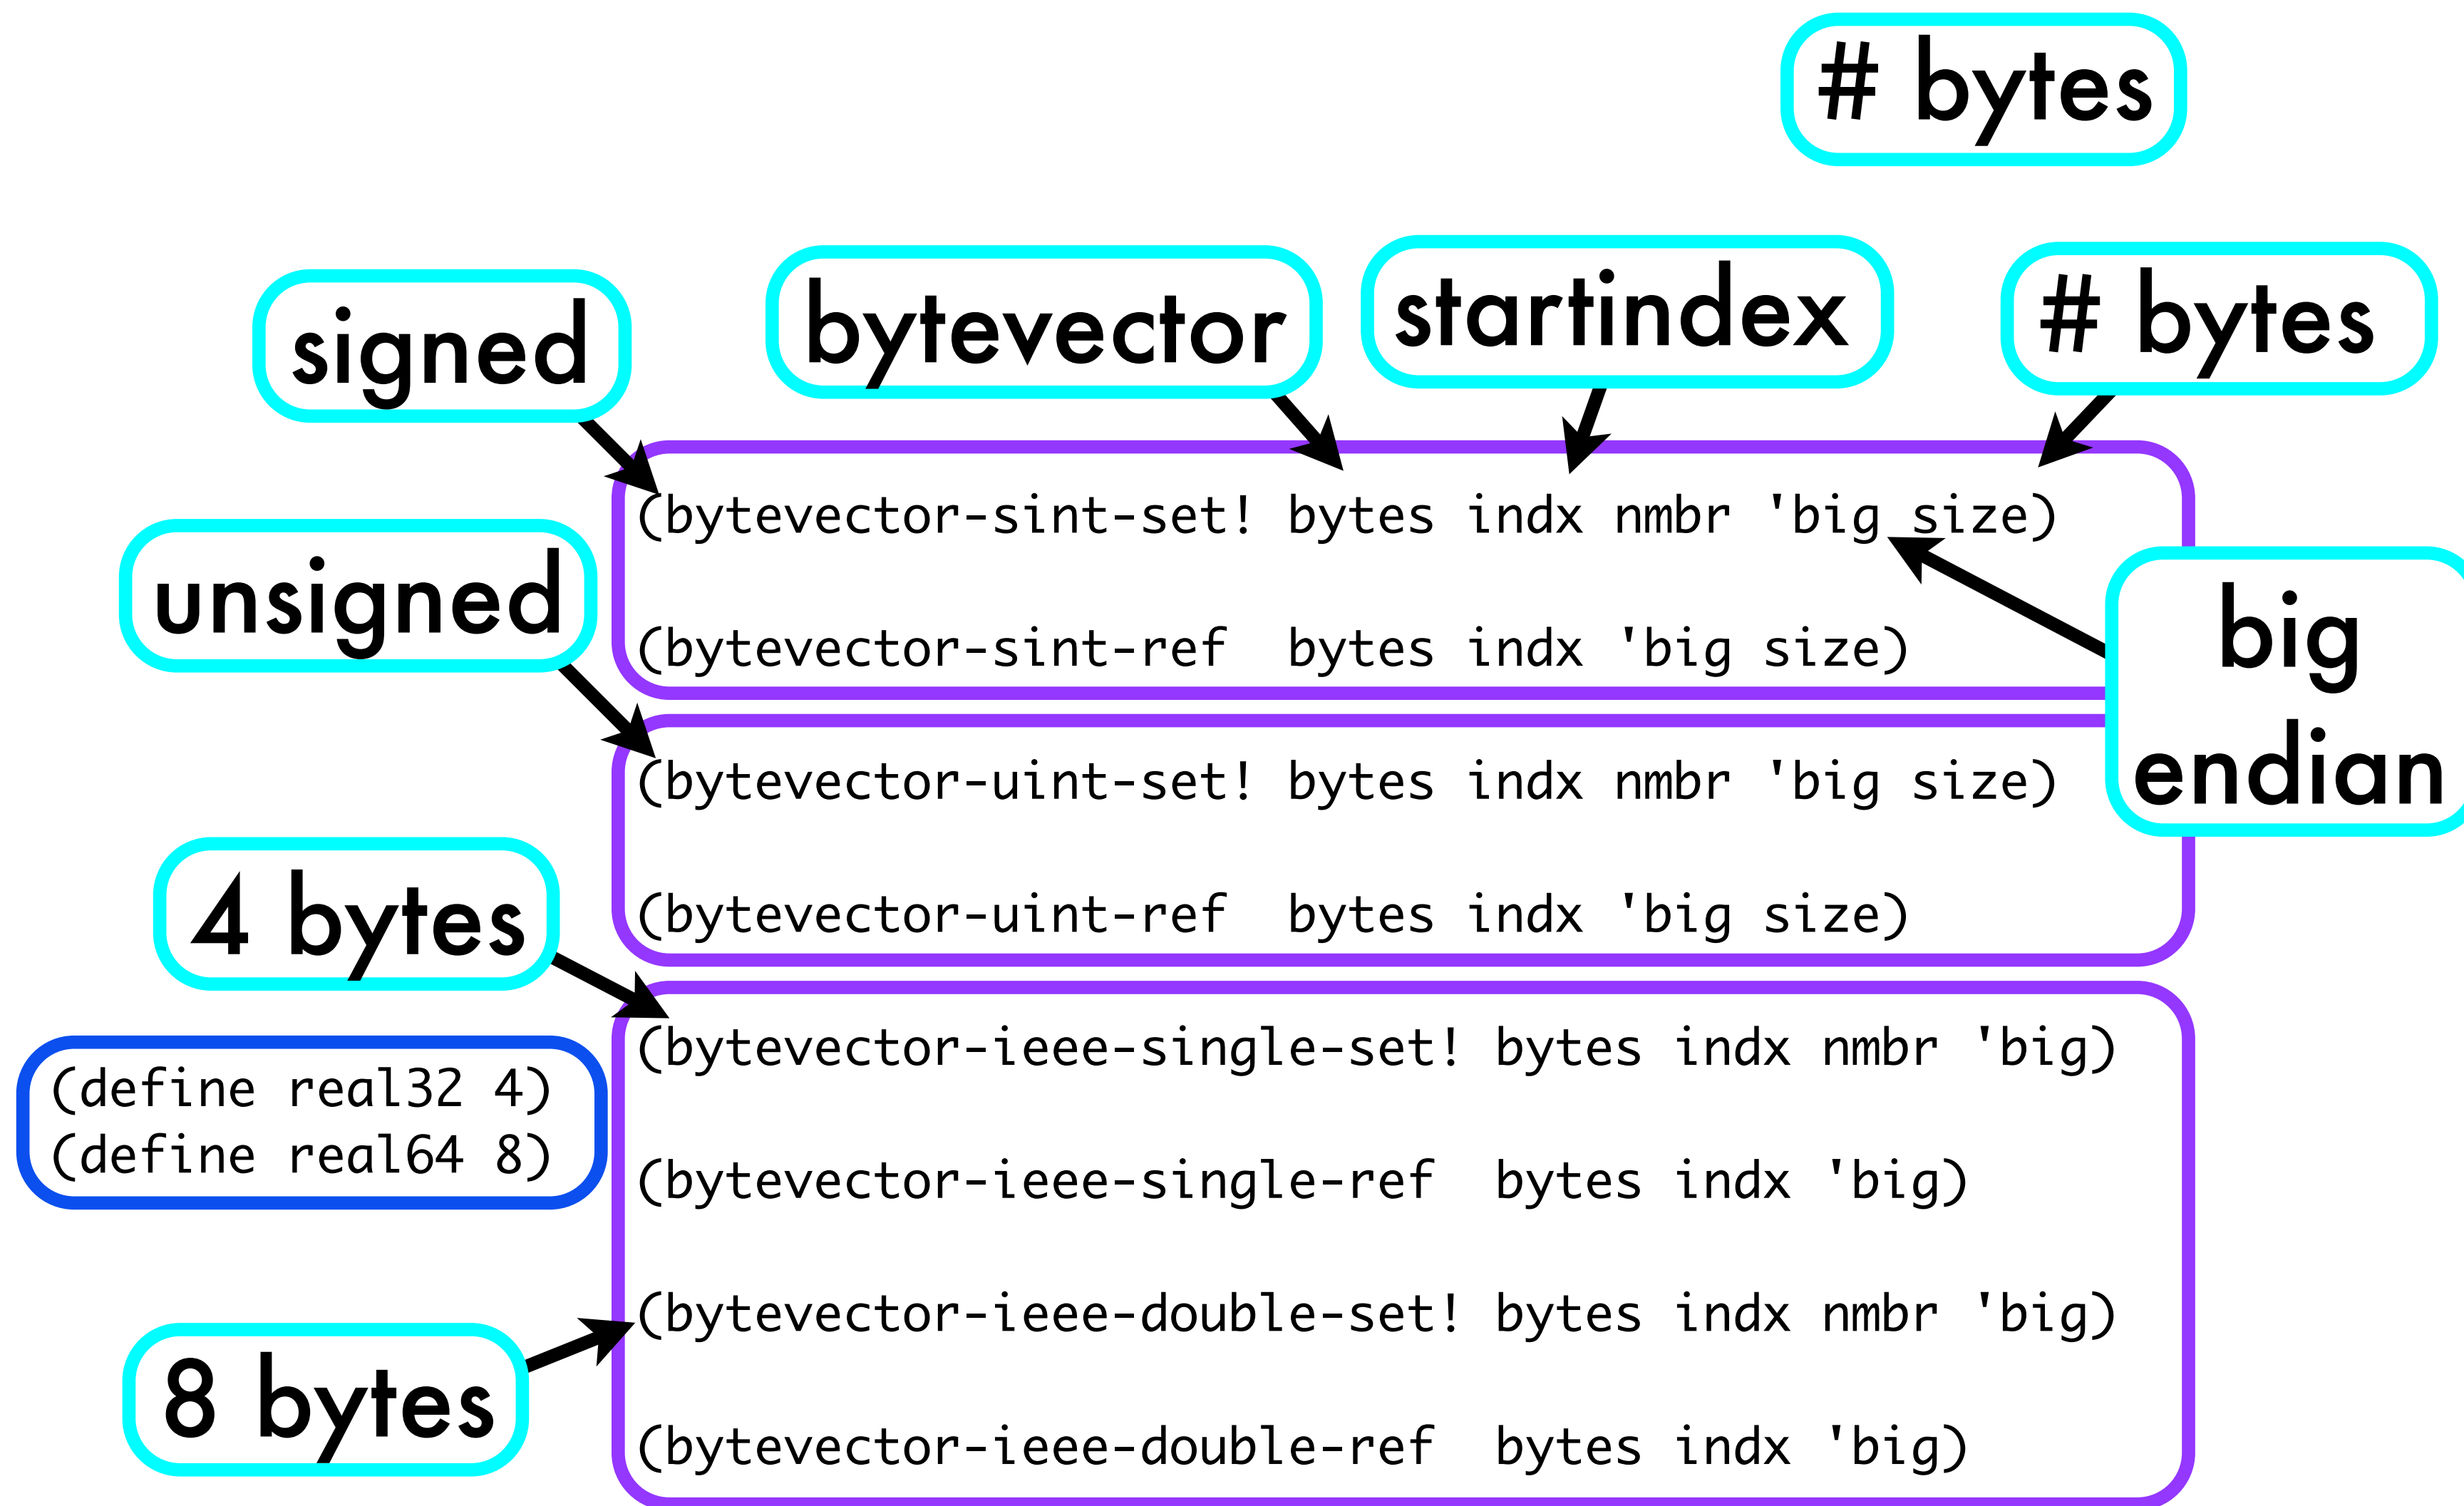
```
> (natural-bytes 0)  
1  
> (natural-bytes 1)  
1  
> (natural-bytes 77)  
1  
> (natural-bytes 255)  
1  
> (natural-bytes 256)  
2  
> (natural-bytes 1234567)  
3
```

```
> (integer-bytes 1)  
1  
> (integer-bytes 127)  
1  
> (integer-bytes 128)  
2  
> (integer-bytes -127)  
1  
> (integer-bytes -128)  
1  
> (integer-bytes -129)  
2
```

```
> (log 123 256)  
0.867814313167405  
> (ceiling (log 123 256))  
1.0  
> (ceiling (log 1234567 256))  
3.0  
> (ceiling (log 1 256))  
0  
> (ceiling (log 0 256))  
log: division by zero
```

(import (scheme inexact))

Getallen in Bytevectors 2/2



Strings in Bytevectoren

`(string->utf8 str)`

`(utf8->string bytes)`

```
> (bytevector-length (string->utf8 "hello"))
5
> (bytevector-length (string->utf8 "Здравствуйте!"))
25
> (bytevector-length (string->utf8 "你好"))
6
```

Later nodig...


```
(define (utf8-sentinel-for nmb-r-byts)
  (define byts (make-bytevector nmb-r-byts))
  (define (fill! offset rem)
    (cond ((= rem 1)
           (bytevector-u8-set! byts offset 127))
          ((= rem 2)
           (bytevector-u8-set! byts offset 223)
           (bytevector-u8-set! byts (+ offset 1) 191))
          (else
           (bytevector-u8-set! byts offset 239)
           (bytevector-u8-set! byts (+ offset 1) 191)
           (bytevector-u8-set! byts (+ offset 2) 191)
           (if (> rem 3)
               (fill! (+ offset 3) (- rem 3))))))
  (fill! 0 nmb-r-byts)
  (utf8->string byts))
```

Iedere lengte
heeft zijn
eigen $+\infty$

```
> (string<? "д" (utf8-sentinel-for 2))
#t
> (string<? "д" (utf8-sentinel-for 1))
#f
```

Waar zitten die bytes?

Centraal Geheugen



Duur

Volatiel

Snel

Periferisch Geheugen



Persistent

Traag


Supergoedkoop

Zeer duur

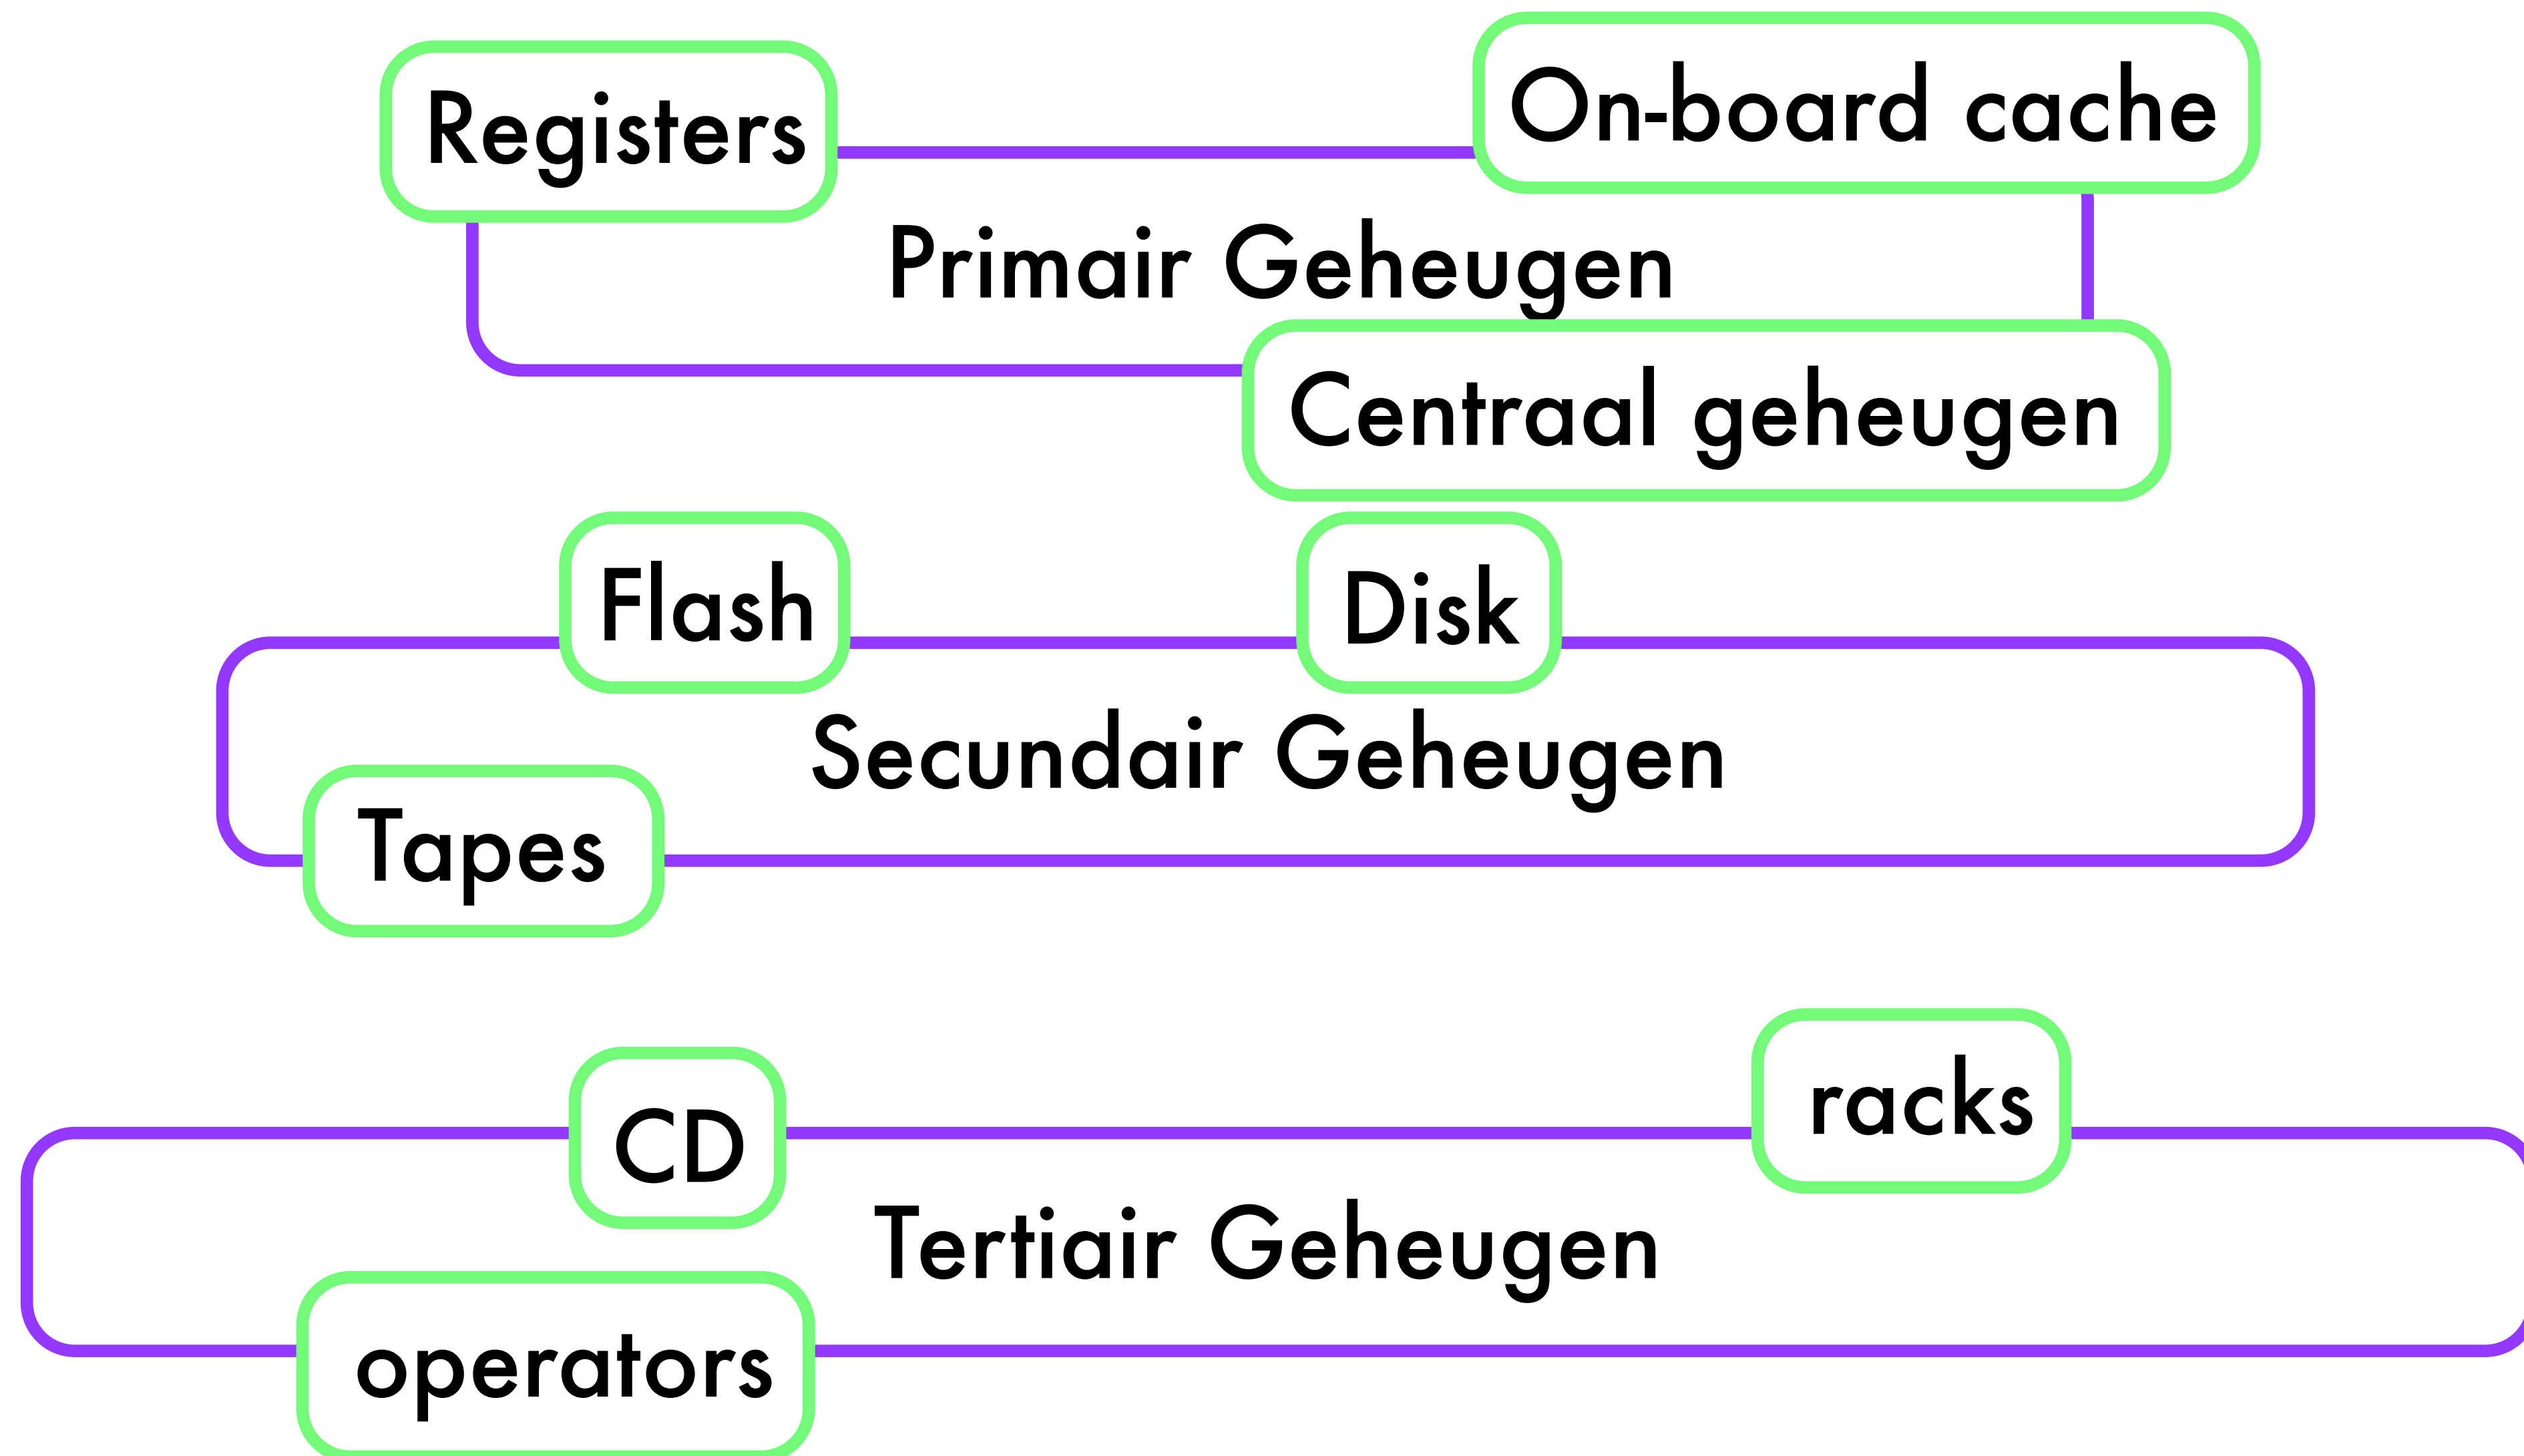
Volatiel

Supersnel

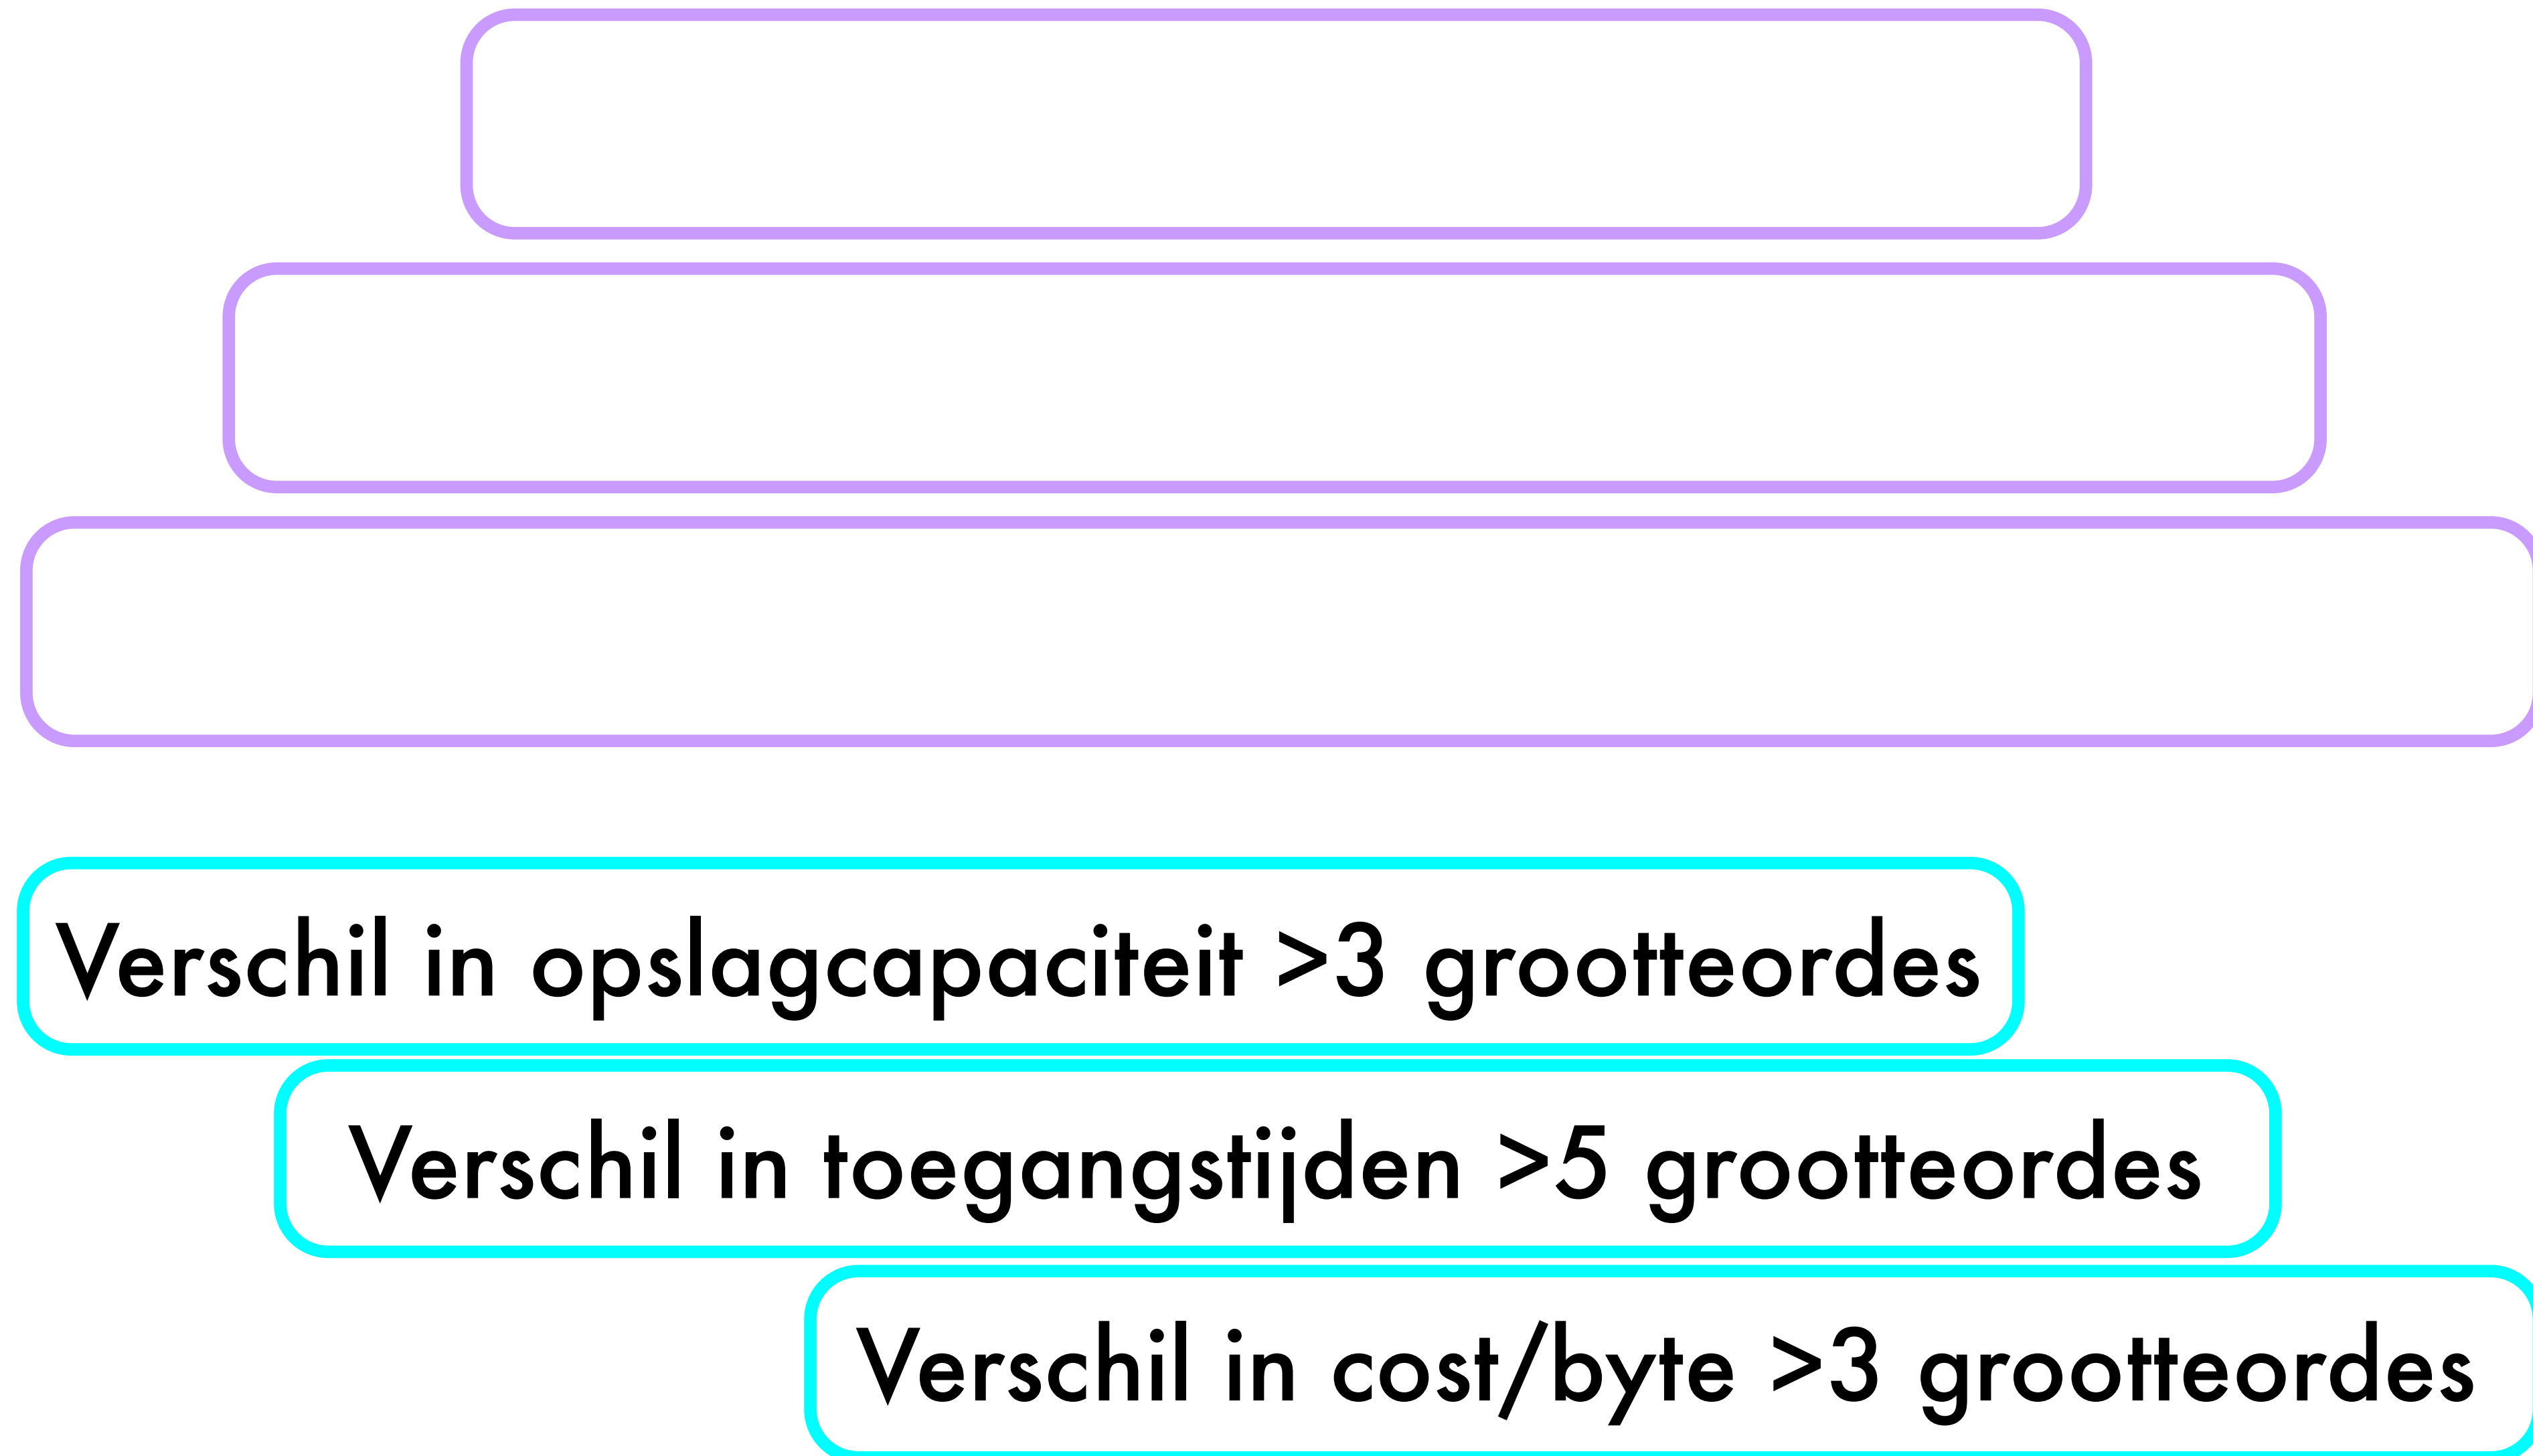
Registers + Caches



De Geheugenhiërarchie




De Geheugenhiërarchie




Disks

Een as met
schijven




Een lees- en
schrijfkop



Magnetiseerbaar
materiaal + hysteresis

Geformateerde Disk




Georganiseerd in tracks en sectoren. Op elke track n sector bevindt zich een blok.

Een blok bevat typisch tussen de 512 bytes en 4K data.

Een blok is de lees- en schrijfeenheid.

Een blok lezen of schrijven



Fase 1: De kop bewegen naar de juiste track.

Fase 2: Wachten tot de juiste sector passeert.

Fase 3: Het blok effectief lezen of schrijven.

$$T = T_{\text{seek}} + T_{\text{latency}} + T_{\text{transfer}}$$

Disk Abstractie

ADT disk

block-size

number

disk-size

number

block-ptr-size

number

block-idx-size

number

new

(string → disk)

mount

(string → disk)

unmount

(disk → ∅)

disk?

(any → boolean)

name

(disk → string)

read-block

(disk number → block)

Een groep blokken met
nummer $\in [0 .. \text{disksize}-1]$

Iedere blok bevat
block-size bytes

```
(define block-size 50)
(define disk-size 200)
(define block-ptr-size (natural-bytes disk-size))
(define block-idx-size (natural-bytes block-size))
```

aantal blocks

size van pointer naar
block binnen disk

size van pointer naar
offset binnen block

Elk bloknummer past in
block-ptr-size bytes

Elke blokindex past in
block-idx-size bytes

Disk (image): Implementatie

```
(define-record-type disk
  (make-disk n)
  disk?
  (n name))
```

Bezet een stuk
van de echte disk

```
(define (new name)
  (define port (open-output-file name #:exists 'truncate))
  (define zeroes (make-bytevector block-size 0))
  (let (low-level-format
        ((block-nr 0))
        (write-bytevector zeroes port)
        (if (< (+ 1 block-nr) disk-size)
            (low-level-format (+ 1 block-nr))))
    (close-port port)
    (make-disk name)))
```

```
(define (mount name)
  (make-disk name))
```

```
(define (unmount disk)
  '())
```

Maakt de
blokken aan

open-output-file, make-bytevector
write-bytevector, close-port

∈ R7RS

Block Abstractie

ADT block

write-block!

(block \rightarrow \emptyset)

block?

(any \rightarrow boolean)

disk

(block \rightarrow disk)

position

(block \rightarrow number)

decode-byte

(block number \rightarrow byte)

encode-byte!

(block number byte \rightarrow \emptyset)

decode-string

(block number number \rightarrow string)

encode-string!

(block number number string \rightarrow \emptyset)

decode-fixed-natural

(block number number \rightarrow natural)

encode-fixed-natural!

(block number number natural \rightarrow \emptyset)

decode-arbitrary-integer

(block number \rightarrow integer \times number)

encode-arbitrary-integer!

(block number integer \rightarrow number)

decode-real

(block number number \rightarrow real)

encode-real!

(block number number real \rightarrow number)

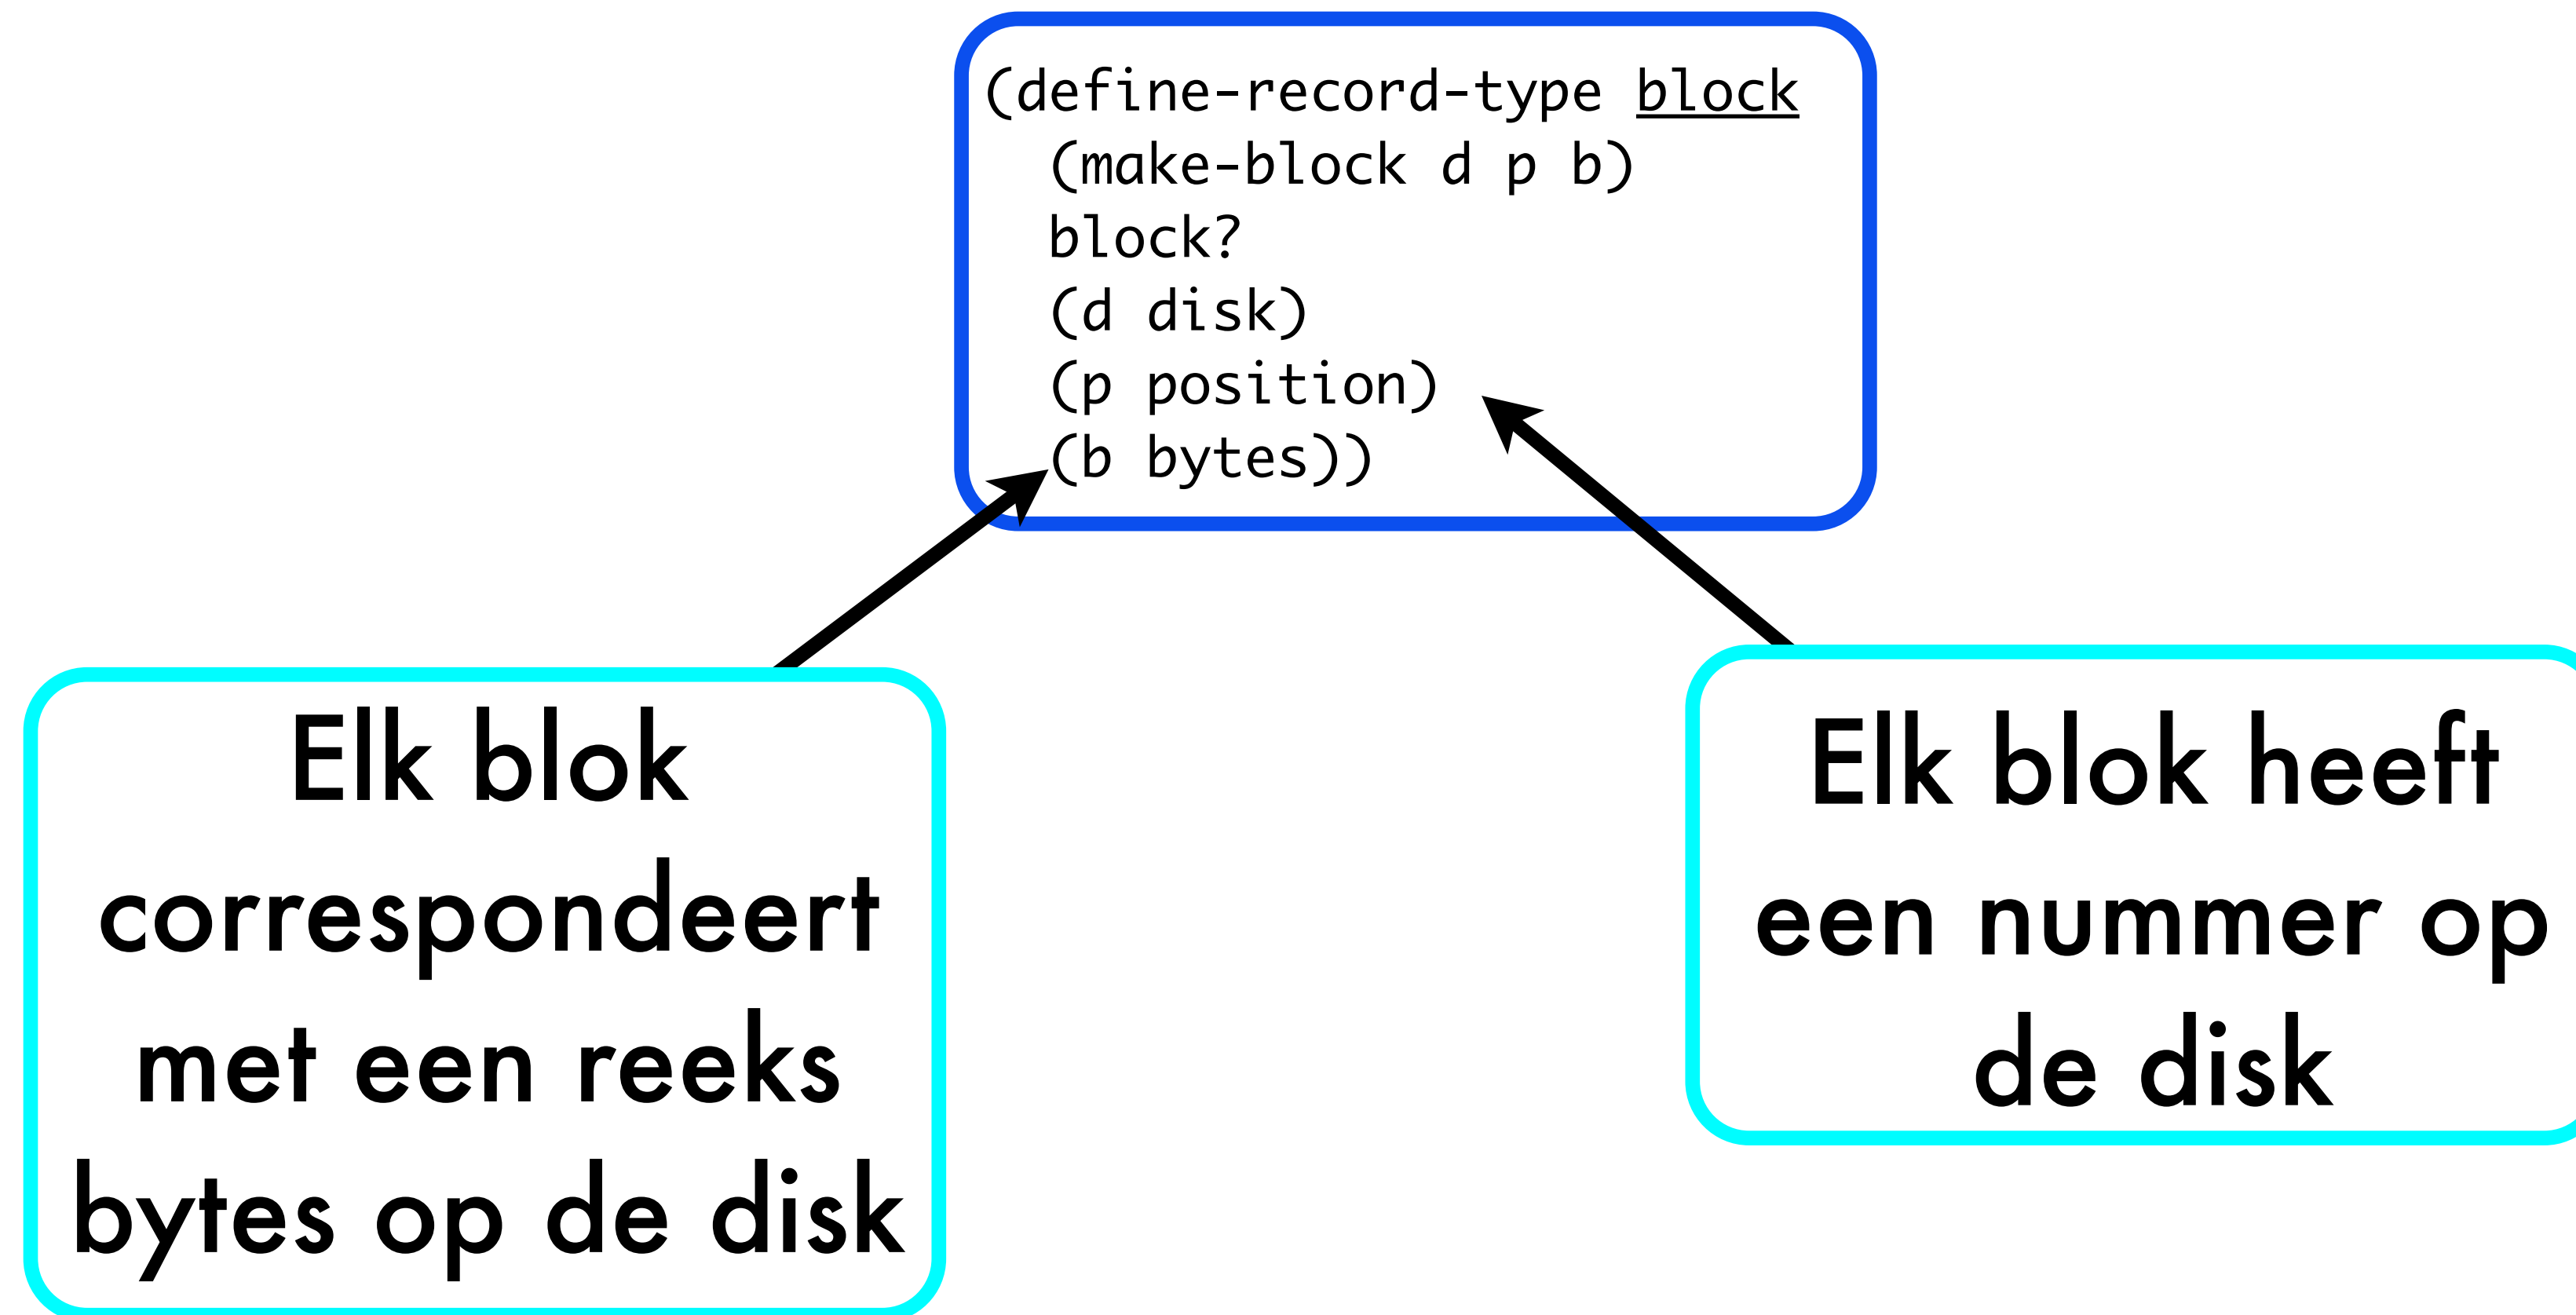
decode-bytes

(block bytevector number number number \rightarrow \emptyset)

encode-bytes!

(block bytevector number number number \rightarrow \emptyset)

Representatie van Blokken



Lezen en Schrijven van een Blok

```
(define (read-block dsk bptr)
  (define port (open-input-file (name dsk)))
  (set-port-position! port (* bptr block-size))
  (let ((byts (read-bytevector block-size port)))
    (close-port port)
    (make-block dsk bptr byts)))
```

```
(define (write-block! blk)
  (define bptr (position blk))
  (define data-byts (bytes blk))
  (define port (open-output-file (name (disk blk)) #:exists 'update))
  (set-port-position! port (* bptr block-size))
  (write-bytevector port data-byts)
  (close-port port))
```

**Zet de "current"
van de output port
op de juiste byte**

**Lees/schrijf de
bytevector horend
bij het blok**

open-input-file, open-output-file,
set-port-position!, read-bytevector,
write-bytevector, close-port

∈ **R7RS**

En/De-coderen van data(1)

```
(define (encode-byte! blk offs byte)
  (bytevector-u8-set! (bytes blk) offs byte))
```

```
(define (decode-byte blk offs)
  (bytevector-u8-ref (bytes blk) offs))
```

offs = offset in het blok

```
(define (encode-fixed-natural! blk offs size nmbr)
  (bytevector-uint-set! (bytes blk) offs nmbr 'big size))
```

```
(define (decode-fixed-natural blk offs size)
  (bytevector-uint-ref (bytes blk) offs 'big size))
```

**Lengte wordt
ook opgeslagen**

```
(define (encode-arbitrary-integer! blk offs nmbr)
  (define size (integer-bytes nmbr))
  (encode-byte! blk offs size)
  (bytevector-sint-set! (bytes blk) (+ offs 1) nmbr 'big size)
  (+ size 1))
```

```
(define (decode-arbitrary-integer blk offs)
  (define size (decode-byte blk offs))
  (define nmbr (bytevector-sint-ref (bytes blk) (+ offs 1) 'big size))
  (cons nmbr (+ offs size 1)))
```


En/De-coderen van data(2)

```
(define (encode-real! blk offs size nmbr)
  (cond ((= size real64)
        (bytevector-ieee-double-set! (bytes blk) offs nmbr 'big)
        size)
        ((= size real32)
        (bytevector-ieee-single-set! (bytes blk) offs nmbr 'big)
        size)
        (else
         (error "illegal real size" size))))

(define (decode-real blk offs size)
  (cond ((= size real64)
        (bytevector-ieee-double-ref (bytes blk) offs 'big))
        ((= size real32)
        (bytevector-ieee-single-ref (bytes blk) offs 'big))
        (else
         (error "illegal real size" size))))
```

```
(define real32 4)
(define real64 8)
```

En/De-coderen van data(3)

```
(define (encode-string! blk off size str)
  (set! str (string->utf8 str))
  (do ((ind 0 (+ ind 1)))
      ((= ind size)
       (let ((byte (if (< ind (bytevector-length str))
                        (bytevector-u8-ref str ind)
                        str-end-byte)))
         (encode-byte! blk (+ off ind) byte))))))

(define (decode-string blk off size)
  (let ((bv (bytevector-u8-set!
              (lambda (bv ind val)
                (bytevector-u8-set! bv ind val)
                bv)))
        (utf8->string
         (let loop
           ((ind 0)
            (if (< ind size)
                (let ((byte (decode-byte blk (+ off ind))))
                  (if (eq? byte str-end-byte)
                      (make-bytevector ind 0)
                      (bytevector-u8-set! (loop (+ ind 1)) ind byte)))
                (make-bytevector ind 0)))))))
    bv))
```


Kopieer alle
bytes en vul
aan met 0'en

Padding

(define str-end-byte 0)

Blijf uitlezen tot
je een 0 ziet

En/De-coderen van rauwe data(4)



```
(define (encode-bytes! blk bytes blk-offs u8-offs size)
  (do ((indx 0 (+ indx 1)))
      ((= indx size)
       (encode-byte! blk
                     (+ blk-offs indx)
                     (bytevector-u8-ref bytes (+ u8-offs indx))))))
```

**Kopieer bytes
in het blk**

```
(define (decode-bytes blk bytes blk-offs u8-offs size)
  (do ((indx 0 (+ indx 1)))
      ((= indx size)
       (bytevector-u8-set! bytes
                           (+ u8-offs indx)
                           (decode-byte blk (+ blk-offs indx)))))
```

Inverse operatie

Het Filestysteem

Unix

Mac OS X

DOS

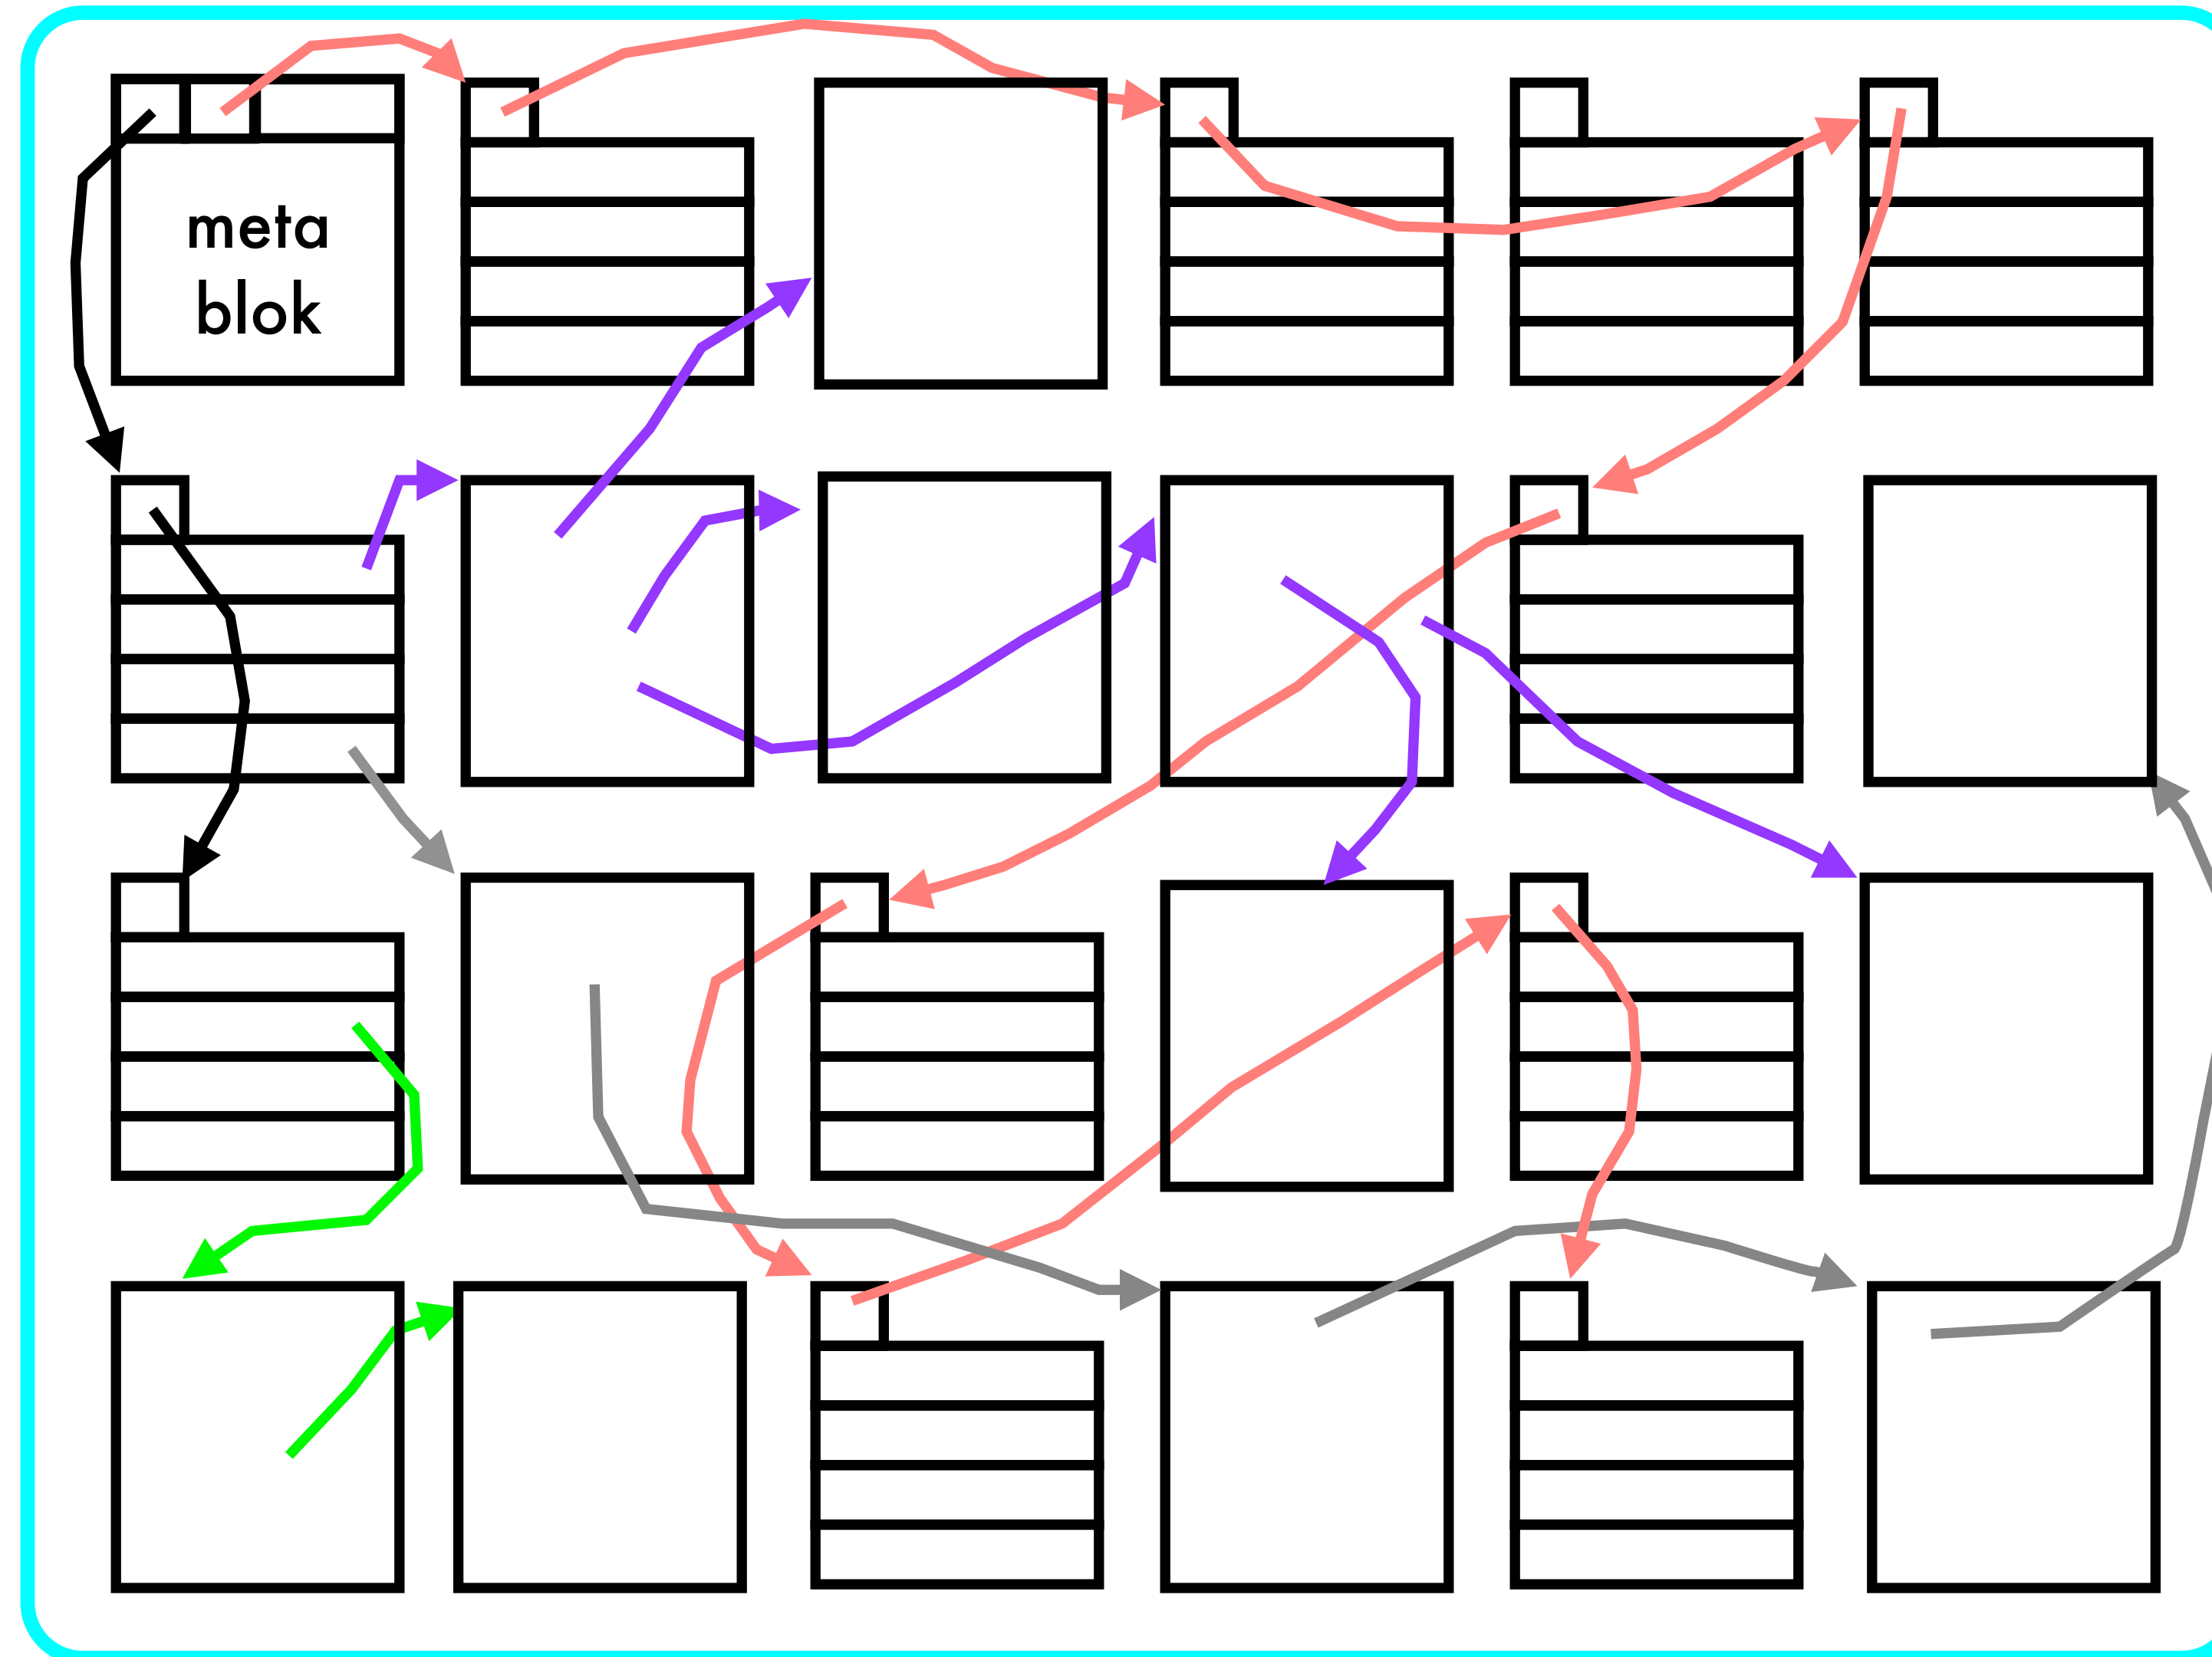
Data wordt georganiseerd in allerlei structuren die ontstaan door blokken aan elkaar te linken. Dit gebeurt door de blokpointer van het ene blok te encoderen als bytes in het andere blok. Een reeks aan elkaar gelinkte blokken noemen we een file.

Sommige blokken behoren tot één of andere file; andere zijn vrij. De vrije blokken worden bijgehouden in een freelist.

Alles begint bij het metablok

De files geven we een naam. De directory is een lijst van filenamen met bloknummer van het begin van de file

Organisatie van de Disk



Filesystem: Constanten

```
(define filename-size 10)
```

```
(define (cap-name name)
  (if (> (string-length name) filename-size)
      (substring name 0 filename-size)
      name))
```

$+\infty$

```
(define sentinel-filename (utf8-sentinel-for filename-size))
```

```
(define null-block 0)
(define meta-bptr null-block)
```

```
(define (null-block? bptr)
  (= bptr null-block))
```

0 is het bloknummer van het metablok en dient dus ook als ()

mag geen onderdeel
van file zijn

```
(define (read-meta-block disk)
  (disk:read-block disk meta-bptr))
```

```
(define (write-meta-block! meta)
  (disk:write-block! meta))
```

Enkel-Linken van Blokken

```
(define next-offset 0)
```

```
(define (next-bptr blk)  
  (disk:decode-fixed-natural blk next-offset disk:block-ptr-size))
```

```
(define (next-bptr! blk bptr)  
  (disk:encode-fixed-natural! blk next-offset disk:block-ptr-size bptr))
```

```
(define (has-next? blk)  
  (not (null-block? (next-bptr blk))))
```

Layout van het Metablok

```
(define directory-offset 0)
(define freelist-offset disk:block-ptr-size)
(define available-offset (* 2 disk:block-ptr-size))


(define (directory meta)
  (disk:decode-fixed-natural meta directory-offset disk:block-ptr-size))
(define (directory! meta blk)
  (disk:encode-fixed-natural! meta directory-offset disk:block-ptr-size blk))

(define (freelist meta)
  (disk:decode-fixed-natural meta freelist-offset disk:block-ptr-size))
(define (freelist! meta flst)
  (disk:encode-fixed-natural! meta freelist-offset disk:block-ptr-size flst))

(define (blocks-free meta)
  (disk:decode-fixed-natural meta available-offset disk:block-ptr-size))
(define (blocks-free! meta free)
  (disk:encode-fixed-natural! meta available-offset disk:block-ptr-size free))
```

Structuur van de Directory

(name,bptr) - slots



Enkel volledig lege blokken worden opnieuw in de freelist gehangen

Layout van één Directoryblok

```
(define slot-size      (+ filename-size disk:block-ptr-size))
(define nr-of-dir-slots (quotient (- disk:block-size disk:block-ptr-size)
                                   slot-size))
```

(name,bptr) - slots

```
(define (dir-name/bptr! blk slot name bptr)
  (define offn (+ disk:block-ptr-size (* slot slot-size)) )
  (define offp (+ disk:block-ptr-size (* slot slot-size) filename-size))
  (disk:encode-string! blk offn filename-size name)
  (disk:encode-fixed-natural! blk offp disk:block-ptr-size bptr))
```

```
(define (dir-name blk slot)
  (define offn (+ disk:block-ptr-size (* slot slot-size)))
  (disk:decode-string blk offn filename-size))
```

```
(define (dir-bptr blk slot)
  (define offp (+ disk:block-ptr-size (* slot slot-size) filename-size))
  (disk:decode-fixed-natural blk offp disk:block-ptr-size))
```

```
(define (empty-slot? blk slot)
  (string=? sentinel-filename (dir-name blk slot)))
```

```
(define (at-end? blk slot)
  (= slot nr-of-dir-slots))
```


High-level Format van de Disk

Het metablok bevat een pointer naar de directory, naar de freelist en houdt het aantal vrije blokken bij

Het metablok zélf is niet vrij

Ieder blok in de freelist bevat het bloknummer van zijn opvolger

```
(define (format! disk)
  (define meta (read-meta-block disk))
  (directory! meta null-block)
  (freelist! meta (+ meta-bptr 1))
  (blocks-free! meta (- disk:disk-size 1))
  (write-meta-block! meta)
  (let high-level-format
    ((bptr (+ meta-bptr 1)))
    (let ((block (disk:read-block disk bptr)))
      (cond ((< (+ bptr 1) disk:disk-size)
              (next-bptr! block (+ bptr 1))
              (disk:write-block! block)
              (high-level-format (+ bptr 1)))
            (else
             (next-bptr! block null-block)
             (disk:write-block! block))))))
  disk)
```

Beheer van Vrije Blokken (1)

Een blok uit de freelist
"poppen": $O(1)$

```
(define (new-block disk)
  (define meta (read-meta-block disk))
  (define flst (freelist meta))
  (if (null-block? flst)
      (error "disk full! (new-block)" disk)
      (let ((blk (disk:read-block disk flst)))
        (blocks-free! meta (- (blocks-free meta) 1))
        (freelist! meta (next-bptr blk))
        (write-meta-block! meta)
        blk))))
```

Een blok alloceren,
initialiseren en vastlinken

```
(define (fresh-block! disk next!)
  (define next (new-block disk))
  (next-bptr! next null-block)
  (do ((slot 0 (+ slot 1)))
      ((at-end? next slot)
       (next! next)
       next)
      (dir-name/bptr! next slot sentinel-filename null-block))))
```

Beheer van Vrije Blokken (2)

Een blok in de freelist
"pushen": $O(1)$

```
(define (delete-block blk)
  (define disk (disk:disk blk))
  (define meta (read-meta-block disk))
  (next-bptr! blk (freelist meta))
  (disk:write-block! blk)
  (freelist! meta (disk:position blk))
  (blocks-free! meta (+ (blocks-free meta) 1))
  (write-meta-block! meta))
```

```
(define (delete-chain! disk bptr)
  (unless (null-block? bptr)
    (let* ((blk (disk:read-block disk bptr))
          (next (next-bptr blk)))
      (delete-block blk)
      (delete-chain! disk next))))
```

Een hele lijst van
blokken vrijgeven

Filesystem Operatie #1: mk

```
(define (mk disk name bptr)
  (define meta (read-meta-block disk))
  (let loop-dir
    ((dptr (directory meta))
     (new! (lambda (newb)
              (let ((meta (read-meta-block disk)))
                (directory! meta (disk:position newb))
                (write-meta-block! meta))))))
    (let ((blk (if (null-block? dptr)
                   (fresh-block! disk new!)
                   (disk:read-block disk dptr))))
      (let loop-block
        ((slot 0)
         (cond ((at-end? blk slot)
                  (loop-dir (next-bptr blk)
                             (lambda (newb)
                               (next-bptr! blk (disk:position newb))
                               (disk:write-block! blk))))
                ((empty-slot? blk slot)
                 (dir-name/bptr! blk slot name bptr)
                 (disk:write-block! blk))
                (else
                 (loop-block (+ slot 1)))))))
        (loop-block (+ slot 1))))))
```

Een slot aanmaken
in de directory

outer loop:
directory blokken

inner loop: slot
zoeken in een blok

Filesystem Operatie #2: rm

```
(define (rm disk name)
  (define meta (read-meta-block disk))
  (set! name (cap-name name))
  (let loop-dir
    ((bptr (directory meta))
     (nxt! (lambda (next)
              (directory! meta next)
              (write-meta-block! meta))))))
  (let ((blk (if (null-block? bptr)
                 (error "file not found (rm)" name)
                 (disk:read-block disk bptr))))
    (let loop-block
      ((slot 0)
       (seen #f))
      (cond ((at-end? blk slot)
              (loop-dir (next-bptr blk) (lambda (next)
                                           (next-bptr! blk next)
                                           (disk:write-block! blk))))
            ((empty-slot? blk slot)
              (loop-block (+ slot 1) seen))
            ((string=? name (dir-name blk slot))
              (dir-name/bptr! blk slot sentinel-filename null-block)
              (disk:write-block! blk)
              (if (not seen)
                  (maybe-delete-block! blk slot nxt!)))
            (else
              (loop-block (+ slot 1) #t))))))
```

Een slot verwijderen
uit de directory

outer loop:
directoryblokken

inner loop: slot
zoeken in een blok

Filesystem Operatie #2: rm (ctd)

```
(define (maybe-delete-block! blk slot next!)  
  (cond  
    ((at-end? blk slot)  
     (next! (next-bptr blk))  
     (delete-block blk))  
    ((empty-slot? blk slot)  
     (maybe-delete-block! blk (+ slot 1) next!))))
```


Operaties #3&4: whereis, ls

```
(define (ls disk)
  (define meta (read-meta-block disk))
  (define bptr (directory meta))
  (if (null-block? bptr)
      ()
      (let traverse-dir
```

**De directory
opsommen als lijst**

```
(define (whereis disk name)
  (define meta (read-meta-block disk))
  (define bptr (directory meta))
  (set! name (cap-name name))
  (if (null-block? bptr)
      0
      (let traverse-dir
        ((blk (disk:read-block disk bptr))
         (slot 0))
        (cond ((at-end? blk slot)
               (if (has-next? blk)
                   (traverse-dir (disk:read-block disk (next-bptr blk)) 0)
                   null-block))
              ((string=? name (dir-name blk slot))
               (dir-bptr blk slot))
              (else
               (traverse-dir blk (+ slot 1))))))))
```

**Een file opzoeken
in de directory**

Sequentiële Files

Populairste
soort file

Lange rijen van Scheme waarden op disk

ADT output-file

```
new
  ( disk string → output-file )
name
  ( output-file → string )
open-write!
  ( disk string → output-file )
close-write!
  ( output-file → ∅ )
reread!
  ( output-file → input-file )
write!
  ( output-file any → ∅ )
delete!
  ( output-file → ∅ )
```

ADT input-file

```
name
  ( input-file → string )
open-read!
  ( disk string → input-file )
rewrite!
  ( input-file → output-file )
close-read!
  ( input-file → ∅ )
has-more?
  ( input-file → boolean )
read
  ( input-file → any )
peek
  ( input-file → any )
delete!
  ( input-file → ∅ )
```

reread! en rewrite!
"spoelen de file terug"

Voorbeeld

```
(define d (disk:new "My Computer"))  
(fs:format! d)  
  
(define f (out:new d "TestFile"))  
  
(out:write! f 3.14)  
(out:write! f 42)  
(out:write! f "Done!")  
  
(out:close-write! f)  
  
(set! f (in:open-read! d "TestFile"))  
(display (in:read f))(newline)  
(display (in:read f))(newline)  
(display (in:read f))(newline)  
  
(in:close-read! f)
```

```
Welcome to DrRacket, version 5.1.3 [3m].  
Language: r6rs [custom].  
3.14  
42  
Done!  
>
```

Duale Representatie

```
(define-record-type sequential-file
  (make d n h b)
  sequential-file?
  (d disk)
  (n name)
  (h header header!)
  (b buffer buffer!))
```

De header bevat een referentie naar het 1ste blok v/d file. De buffer is een kopie in centraal geheugen van één blok van de file. Slechts 2 blokken in het geheugen!

Header

```
(define frst-offs 0)
(define curr-offs disk:block-ptr-size)

(define (first hder)
  (disk:decode-fixed-natural hder frst-offs disk:block-ptr-size))
(define (first! hder bptr)
  (disk:encode-fixed-natural! hder frst-offs disk:block-ptr-size bptr))
(define (current hder)
  (disk:decode-fixed-natural hder curr-offs disk:block-idx-size))
(define (current! hder offs)
  (disk:encode-fixed-natural! hder curr-offs disk:block-idx-size offs))
```

Creëren / Openen

```
(define (new disk name)
  (define hder (fs:new-block disk))
  (define bffr (fs:new-block disk))
  (define file (make disk name hder bffr))
  (fs:mk disk name (disk:position hder))
  (first! hder (disk:position bffr))
  (fs:next-bptr! bffr fs:null-block)
  (current! hder disk:block-ptr-size)
  (disk:write-block! hder)
  file)
```

**Header maken
& registreren in
de directory**

**Header
opzoeken
in de
directory**

```
(define (open-write! disk name)
  (define bptr (fs:whereis disk name))
  (define hder (disk:read-block disk bptr))
  (define fptr (first hder))
  (define bffr (disk:read-block disk fptr))
  (define file (make disk name hder bffr))
  (current! hder disk:block-ptr-size)
  file)
```

**Current: next
pointer overslaan**

```
(define (open-read! disk name)
  (define bptr (fs:whereis disk name))
  (define hder (disk:read-block disk bptr))
  (define fptr (first hder))
  (define bffr (disk:read-block disk fptr))
  (define file (make disk name hder bffr))
  (current! hder disk:block-ptr-size)
  file)
```

Sluiten / Deleten

```
(define (close-read! file)
  ()); nothing to write
```

```
(define (block-bytes-free file)
  (define hder (header file))
  (define curr (current hder))
  (- disk:block-size curr))
```

```
(define (close-write! file)
  (define hder (header file))
  (define bffr (buffer file))
  (define fdsk (disk file))
  (disk:write-block! hder)
  (if (> (block-bytes-free file) 0)
      (disk:encode-byte! bffr (current hder) eof-tag))
  (let ((rest-list (fs:next-bptr bffr)))
    (fs:next-bptr! bffr fs:null-block)
    (fs:delete-chain! fdsk rest-list))
  (disk:write-block! bffr))
```

**Sluit eventueel de buffer
netjes af, schrijf weg en ruim
ongebruikte blokken op**

```
(define (delete! file)
  (define fnam (name file))
  (define hder (header file))
  (define fdsk (disk file))
  (fs:delete-chain! fdsk (first hder))
  (fs:delete-block hder)
  (fs:rm fdsk fnam))
```

Terugspoelen

```
(define (reread! file)
  (define dsk (disk file))
  (define hder (header file))
  (define fptr (first hder))
  (close-write! file)
  (buffer! file (disk:read-block dsk fptr))
  (current! hder disk:block-pointer-size)
  file)
```

**Current: next
pointer overslaan**

```
(define (rewrite! file)
  (define fdsk (disk file))
  (define hder (header file))
  (define fptr (first hder))
  (close-read! file)
  (buffer! file (disk:read-block fdsk fptr))
  (current! hder disk:block-ptr-size)
  file)
```

Schrijven

Encodeer het type van
de waarde in 1 byte

```
(define (write! file sval)
  (cond ((natural? sval)
        (write-type-tag file natural-tag)
        (write-natural file (exact sval)))
        ((integer? sval)
        (write-type-tag file integer-tag)
        (write-integer file (exact sval)))
        ((real? sval)
        (write-type-tag file decimal-tag)
        (write-real file sval))
        ((string? sval)
        (write-type-tag file string-tag)
        (write-string file sval))
        (error "unsupported type (write!)" sval))))
```

Schrijf de waarde
achter die byte

```
(define natural-tag 0)
(define integer-tag 1)
(define decimal-tag 2)
(define string-tag 3)
(define eof-tag 255)
(define eob-tag 254)
```


Schrijven van ≠ datatypes

```
(define (write-type-tag file ttag)
  (claim-bytes! file 1)
  (let* ((bfr (buffer file))
         (hder (header file))
         (curr (current hder)))
    (disk:encode-byte! bfr curr ttag)
    (current! hder (+ curr 1))))
```

Het type

```
(define (write-natural file nmbr)
  (claim-bytes! file (+ (disk:natural-bytes nmbr) 1))
  (let* ((hder (header file))
```

```
(define (write-integer file nmbr)
  (claim-bytes! file (+ (disk:integer-bytes nmbr) 1))
  (let* ((hder (header file))
```

```
(define (write-real file nmbr)
  (claim-bytes! file disk:real64)
  (let* ((hder (header file))
         (curr (current hder)))
    (disk:encode-byte! bfr curr ttag)
    (current! hder (+ curr 1))))
```

```
(define (write-string file strg)
  (claim-bytes! file 1)
  (let* ((hder (header file))
         (curr (current hder))
         (bfr (buffer file))
         (byts (string->utf8 strg)))
    (disk:encode-byte! bfr curr (bytevector-length byts))
    (current! hder (+ curr 1))
    (rollout-bytes file byts 0)))
```

1 procedure
per type

Garanderen van Plaats

```
(define (claim-bytes! file nmbr)
  (define bffr (buffer file))
  (define hder (header file))
  (define curr (current hder))
  (when (< (block-bytes-free file) nmbr)
    (if (not (= curr disk:block-size))
        (disk:encode-byte! bffr curr eob-tag))
    (provide-next-block! file)))
```

**Ofwel een nieuw
blok; Ofwel een
oud blok**

```
(define (provide-next-block! file)
  (define fdsk (disk file))
  (define hder (header file))
  (define bffr (buffer file))
  (define next (cond ((fs:null-block? (fs:next-bptr bffr))
                      (let ((newb (fs:new-block fdsk)))
                        (fs:next-bptr! newb fs:null-block)
                        (fs:next-bptr! bffr (disk:position newb))
                        (disk:write-block! bffr)
                        newb))
                     (else
                      (disk:write-block! bffr)
                      (disk:read-block fdsk (fs:next-bptr bffr))))))
  (buffer! file next)
  (current! hder disk:block-ptr-size))
```


Lezen

```
(define (read file)
  (define hder (header file))
  (define curr (current hder))
  (let* ((ttag (read-type-tag file))
        (peek file)
        (curr! (current! hder curr-offs))
        (buffer! (buffer! file bffr))
        (res (car (cond ((= ttag natural-tag)
                        (peek-natural file))
                        ((= ttag integer-tag)
                        (peek-integer file))
                        ((= ttag decimal-tag)
                        (peek-real file))
                        ((= ttag string-tag)
                        (peek-string file))
                        ((= ttag eof-tag)
                        (cons () curr-offs))
                        (else
                         (error "unsupported type on file (peek)" ttag)))))))
    (current! hder curr-offs)
    (buffer! file bffr)
    res))
```

Lees de waarde na
het type gelezen te
hebben

peek reset
de toestand

Garanderen van Bytes

```
(define (supply-bytes! file)
  (define bffr (buffer file))
  (define hder (header file))
  (define curr (current hder))
  (if (not (more-on-buffer? file))
      (read-next-block! file)))
```

```
(define (more-on-buffer? file)
  (define bffr (buffer file))
  (define hder (header file))
  (define offs (current hder))
  (and (< offs disk:block-size)
       (not (or (= (disk:decode-byte bffr offs) eob-tag)
                 (= (disk:decode-byte bffr offs) eof-tag)))))
```

```
(define (read-next-block! file)
  (define fdsk (disk file))
  (define hder (header file))
  (define bffr (buffer file))
  (define next-bptr (fs:next-bptr bffr))
  (define next-blck (disk:read-block fdsk next-bptr))
  (buffer! file next-blck)
  (current! hder disk:block-ptr-size))
```

H9-H11: Hét Centrale Thema

Minimaliseren van het aantal bloktransferten

Caching

Een cache is een mechanisme dat twee verschillende geheugens doorverbindt en het verschil in performantiegedrag van die twee geheugens opvangt via statistische technieken.


Een stuk gereserveerd geheugen in het snelle geheugen dat een kopie van een deel van het trage geheugen bevat.

Enkele bytevectoren
in de processor

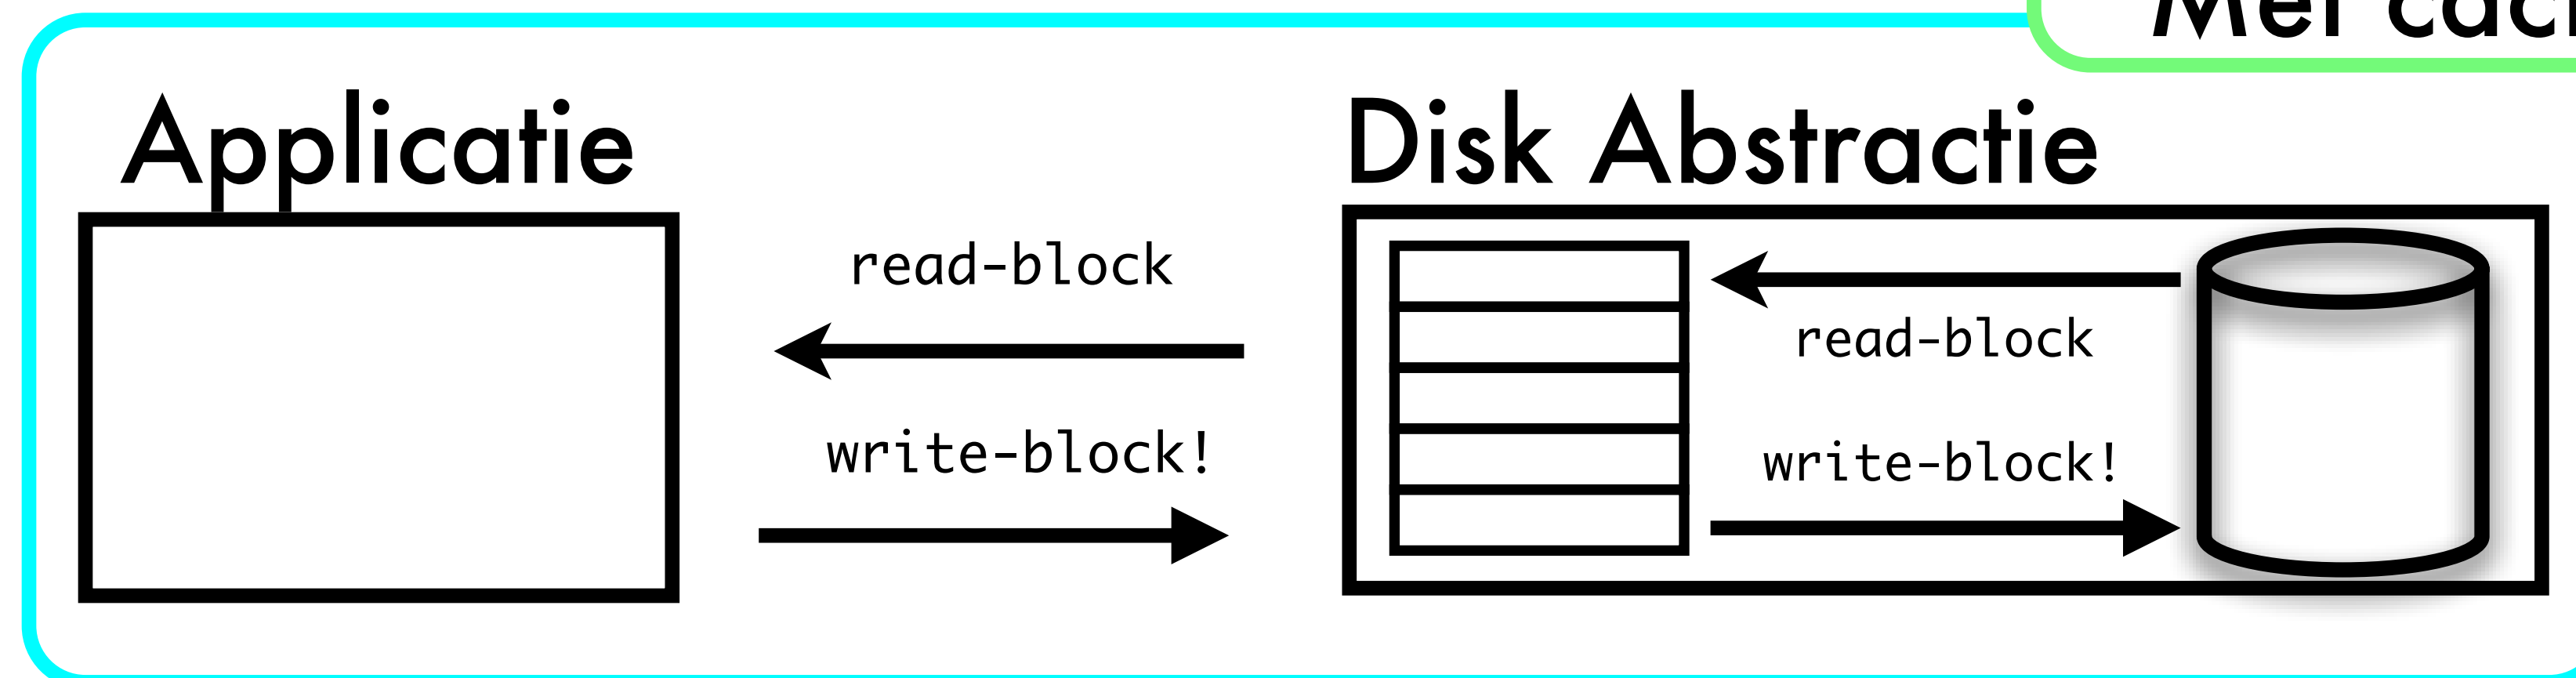
Enkele blokken in
centraal geheugen

Disk Caches

Zonder cache



Met cache



Basis voor Caching: Locality

Temporal Locality: De meest recent gebruikte geheugenlocaties hebben de grootste kans om in de nabije toekomst opnieuw gebruikt te worden.

Spatial Locality: Geheugenlocaties dichtbij de meest recent gebruikte geheugenlocaties hebben de grootste kans om in de nabije toekomst gebruikt te worden.

Selection Sort

Quicksort

Caching Terminologie

Als een benodigde datawaarde tijdens het lezen van de cache gevonden wordt, spreken we van een cache hit. Anders van een cache miss en is een evict nodig alvorens een nieuwe datawaarde in de cache kan geladen worden.

Als een datawaarde in de cache geschreven wordt moet het trage geheugen ook ooit geupdate worden. In een write-through cache gebeurt dat meteen. In een write-back cache gebeurt dat bij een evict. Een dirty-bit is dan nodig voor boekhouding.

Welke Datawaarden Evicten?

Random: Gooi een willekeurige datawaarde uit de cache.

$O(1)$

First-In, First-Out (FIFO): De oudste datawaarde wordt eruit gegooid.

Least Frequently Used (LFU): Gooi de minst gebruikte datawaarde eruit.

$O(\text{size})$

Least Recently Used (LRU): Gooi de datawaarde eruit die het langst niet meer gebruikt is.

Gecachete Disk Blokken

**write-block! maakt
het blok ongelocked**

**read-block maakt
het blok gelocked**

**Elk gebruikt (encode of decode)
actualiseert de time-stamp**

**Encoders maken
het blok "dirty"**

**Eens ge-evict uit de cache,
is het blok niet meer valid**

```
dirty?  
  ( block → boolean )  
dirty!  
  ( block boolean → ∅ )  
locked?  
  ( block → boolean )  
locked!  
  ( block boolean → ∅ )  
time-stamp  
  ( block → time )  
time-stamp!  
  ( block time → ∅ )  
invalidate!  
  ( block → ∅ )  
valid?  
  ( block → boolean )
```

Gecachte Blokken

```
(define-record-type cblock
  (make d l t i b)
  block?
  (d dirty? dirty-set!)
  (l locked? locked!)
  (t time-stamp time-stamp!)
  (i disk)
  (b block block!))

(define (make-block cdsk blk)
  (make #f #t (current-time) cdsk blk))
```

dirty: werd het blok door encoding aangerakt?

locked: werd het block door de applicatie nog niet logisch weggeschreven?

```
(define (dirty! blk)
  (dirty-set! blk #t))

(define (invalidate! blk)
  (block! blk ()))

(define (valid? blk)
  (not (null? (block blk))))
```


Gecachte Encoders

```
(define (make-cached-encoder proc)
  (lambda args
    (define blk (car args))
    (if (not (valid? blk))
        (error "invalidated cblock(cached version of encoder)" blk))
    (time-stamp! blk (current-time))
    (dirty! blk)
    (apply proc (cons (block blk) (cdr args))))))

(define encode-byte!
  (make-cached-encoder disk:encode-byte!))
(define encode-fixed-natural!
  (make-cached-encoder disk:encode-fixed-natural!))
(define encode-arbitrary-integer!
  (make-cached-encoder disk:encode-arbitrary-integer!))
(define encode-real!
  (make-cached-encoder disk:encode-real!))
(define encode-string!
  (make-cached-encoder disk:encode-string!))
(define encode-integer!
  (make-cached-encoder disk:encode-integer!))
```

**De dirty flag
wordt gezet**

**We maken het
blok "jonger"**

**Oorspronkelijke
operatie wordt
opgeroepen**

**Operaties op invalid blokken
worden tegengehouden**

Gecachte Decoders

```
(define (make-cached-decoder proc)
  (lambda args
    (define blk (car args))
    (if (not (valid? blk))
        (error "invalidated cblock(cached version of decoder)" blk))
    (time-stamp! blk (current-time))
    (apply proc (cons (block blk) (cdr args))))))

(define decode-byte
  (make-cached-decoder disk:decode-byte))
(define decode-fixed-natural
  (make-cached-decoder disk:decode-fixed-natural))
(define decode-arbitrary-integer
  (make-cached-decoder disk:decode-arbitrary-integer))
(define decode-real
  (make-cached-decoder disk:decode-real))
(define decode-string
  (make-cached-decoder disk:decode-string))
(define decode-bytes
  (make-cached-decoder disk:decode-bytes))
```

Zelfde idee

Gecachete Disk

ADT disk

block-size

number

disk-size

number

block-ptr-size

number

block-idx-size

number

new

(string → disk)

mount

(string → disk)

unmount

(disk → ∅)

disk?

(any → boolean)

name

(file → string)

size

(file → number)

read-block

(disk number → block)

**ADT blijft
onveranderd**

**Implementatie
bevat echter een
cache in het
centraal geheugen**

De Eigenlijke Cache

```
(define cache-size 10)

(define (cache:new)
  (make-vector cache-size ()))

(define (cache:get cche indx)
  (vector-ref cche indx))

(define (cache:put! cche indx blk)
  (vector-set! cche indx blk))
```

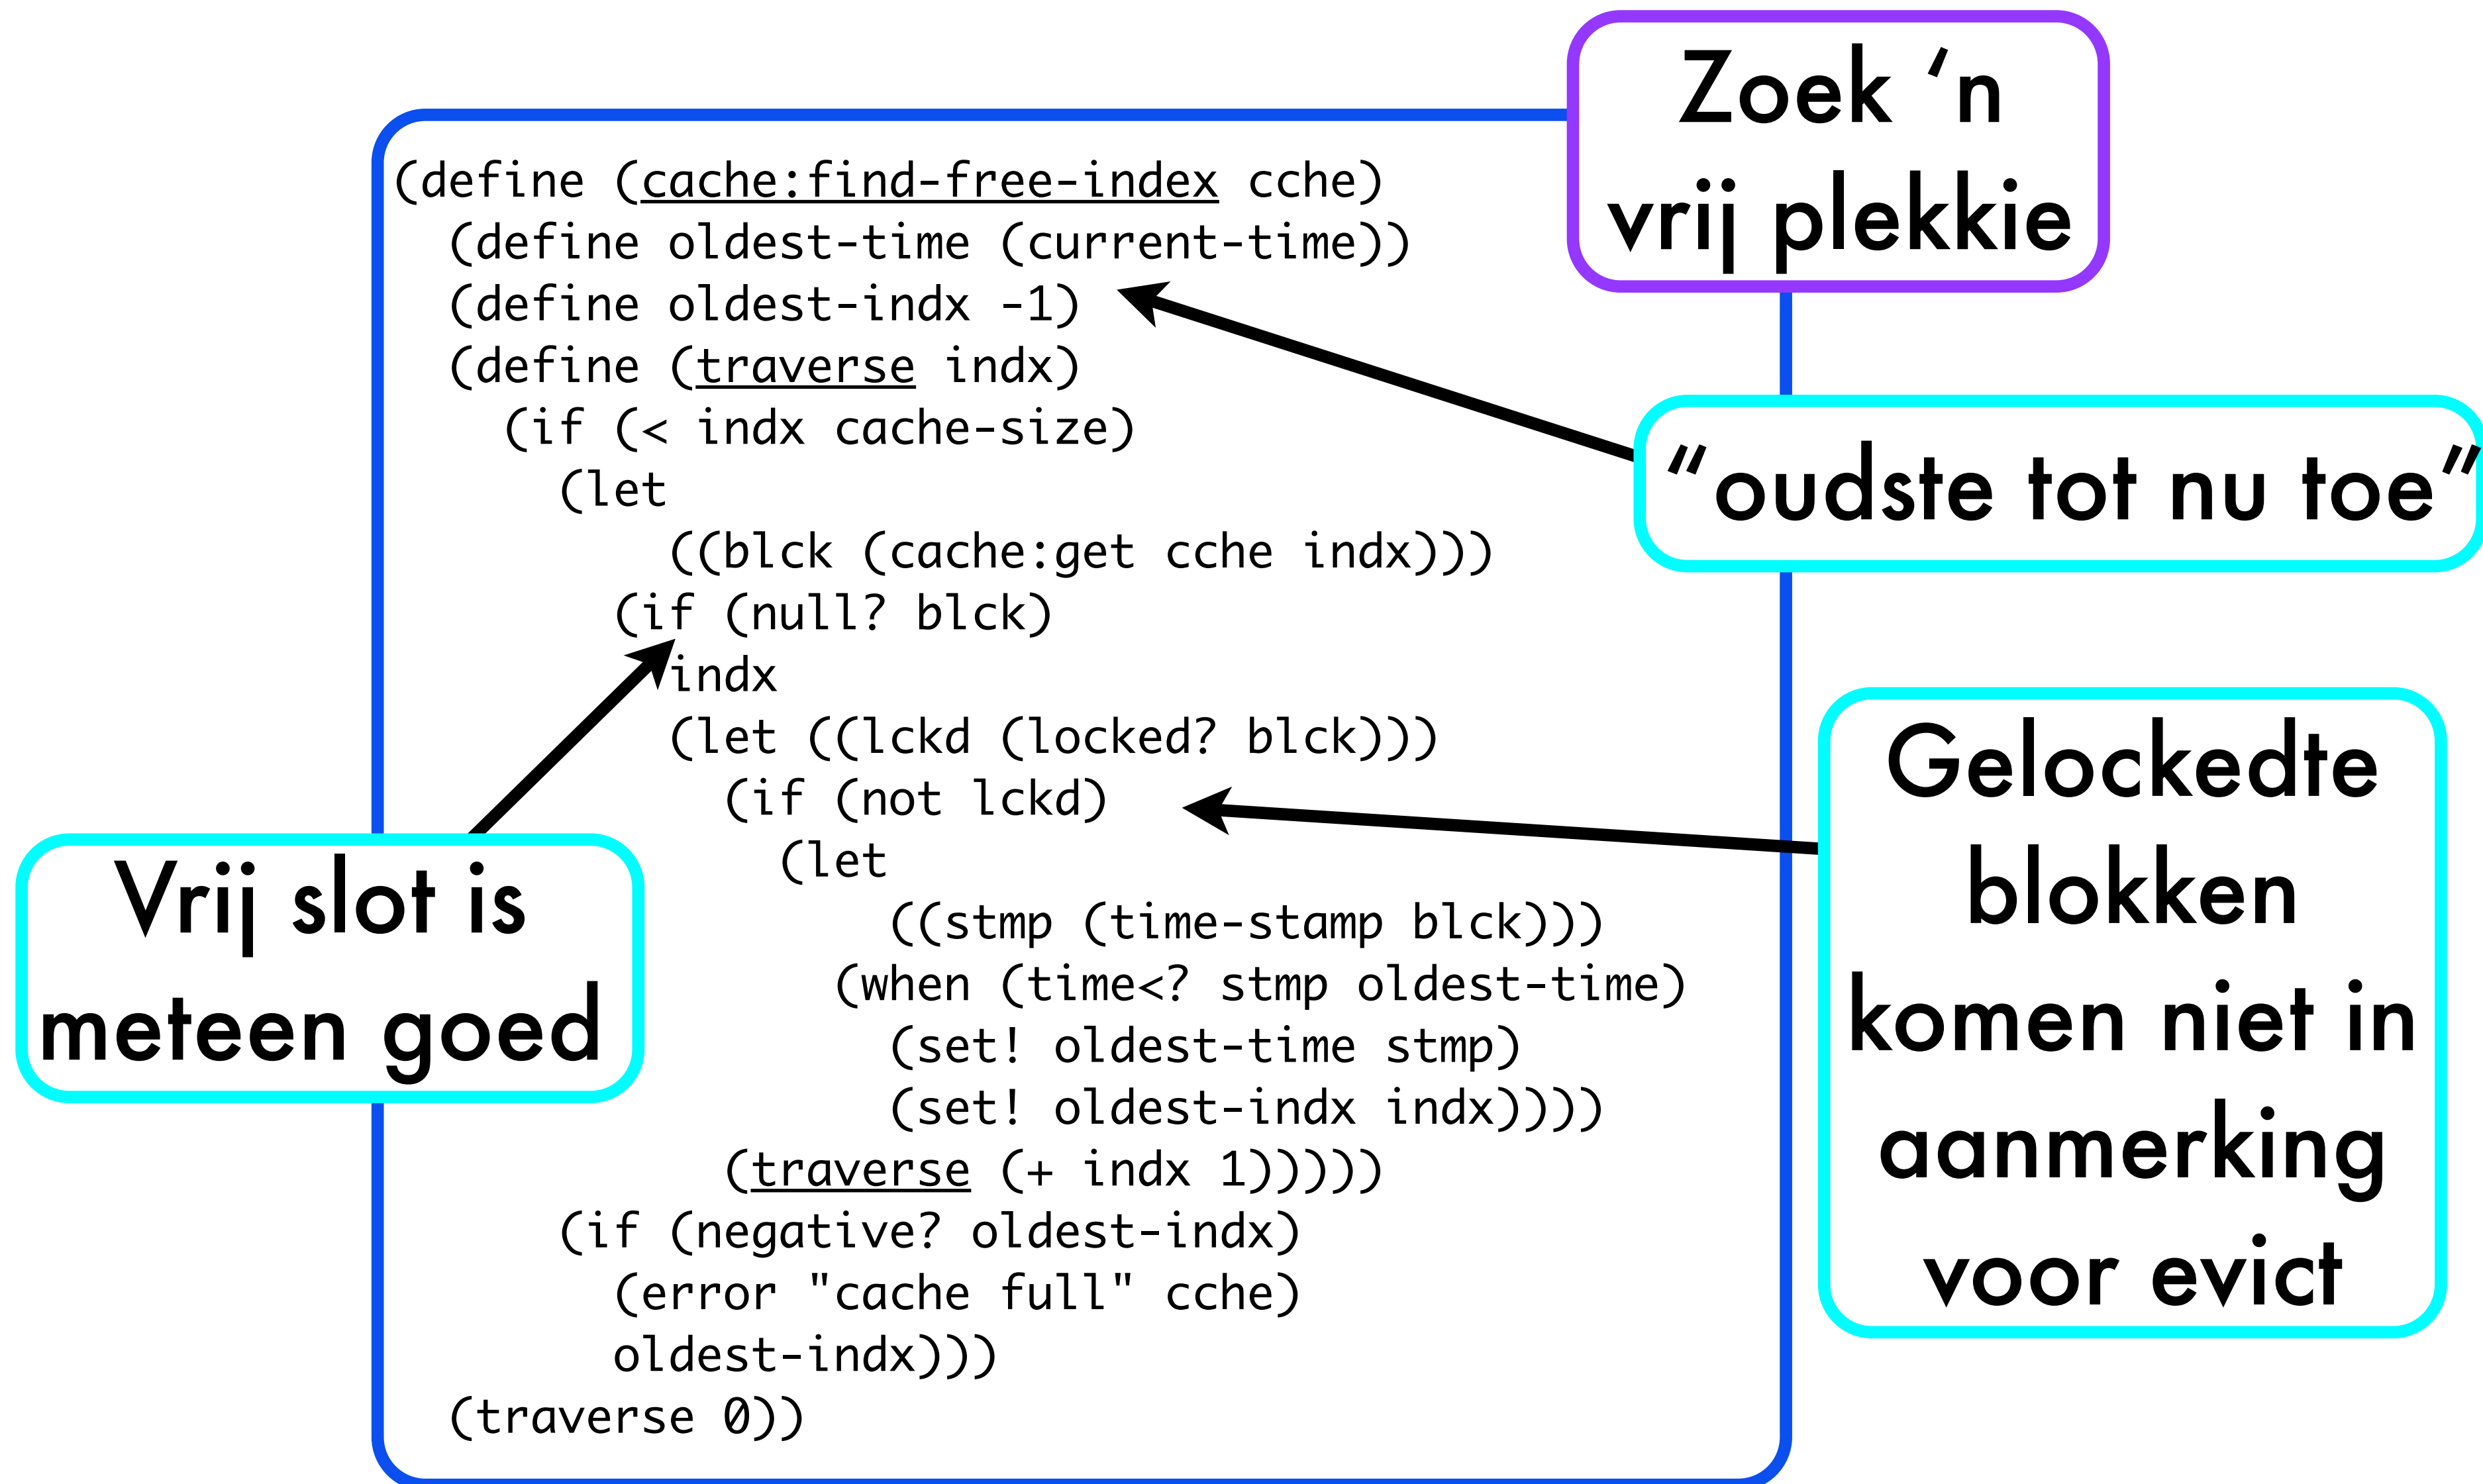
**Gewoon een
vector van blokken**

**Zoek een gegeven blok
nummer op in de cache**

```
(define (cache:find-block cche bptr)
  (define (position-matches? blk)
    (and (not (null? blk))
         (= (position blk) bptr)))
  (let traverse
    ((indx 0)
     (blk (cache:get cche 0)))
    (cond ((position-matches? blk)
           blk)
          ((< (+ indx 1) cache-size)
           (traverse (+ indx 1)
                       (cache:get cche (+ indx 1))))
          (else
           ())))))
```

**Deze cache is fully
associative: ieder blok
kan gelijk waar terecht
komen in de cache.**

De eigenlijke Cache (ctd)



Creëren, Mounten & Unmounten

```
(define-record-type cdisk  
  (make-cdisk v d)  
  disk?  
  (v disk-cache)  
  (d real-disk))
```

**Gecachete disk
= disk + cache**

```
(define (new name)  
  (make-cdisk (cache:new) (disk:new name)))  
  
(define (mount name)  
  (make-cdisk (cache:new) (disk:mount name)))  
  
(define (name cdsk)  
  (disk:name (real-disk cdsk)))
```

**Alle dirty blokken
moeten geschreven
worden**

**Wandel de
volledige
cache af**

```
(define (unmount disk)  
  (define vctr (cache:vector disk))  
  (define (traverse indx)  
    (if (< indx cache-size)  
      (let ((blk (cache:get vctr indx)))  
        (cond  
          ((not (null? blk))  
           (if (dirty? blk)  
               (disk:write-block! (block blk))  
               (invalidate! blk)))  
          (traverse (+ indx 1))))))  
  (traverse 0)  
  (disk:unmount disk))
```

Lezen en Schrijven van Blokken

```
(define (read-block cdsk bptr)
  (define cche (disk-cache cdsk))
  (define blk (cache:find-block cche bptr))
  (if (null? blk)
      (let*
        ((indx (cache:find-free-index cche))
         (blk (cache:get cche indx)))
        (when (not (null? blk))
          (if (dirty? blk)
              (disk:write-block! (block blk))
              (invalidate! blk))
          (set! blk (make-block cdsk (disk:read-block (real-disk cdsk) bptr)))
          (cache:put! cche indx blk)))
      (locked! blk #t)
      blk)
```

Indien niet in cache,
lees van de disk,
eventueel na een evict

Gelezen blokken
worden gelocket

```
(define (write-block! blk)
  (if (not (valid? blk))
      (error "invalidated block(write-block!)" blk))
  (locked! blk #f))
```

write schrijft het block
niet meteen weg!

en worden
door schrijven
weer unlocked

Hoofdstuk 14

14.1 Wat is Data? Bits, bytes, bytevectoren

14.2 De disk abstractie

14.3 Het filesystem

14.4 Sequentiële Files

14.5 Caching

